

# Probabilistic Linear Solvers

Jonathan Wenger

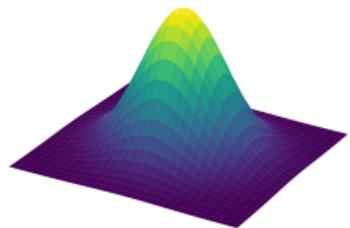
COLUMBIA | Zuckerman Institute



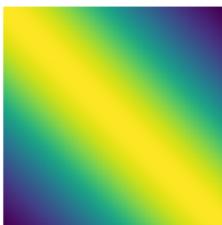
# Linear Systems are Everywhere in Scientific Computing

Arguably, the most fundamental numerical task in scientific computing and machine learning.

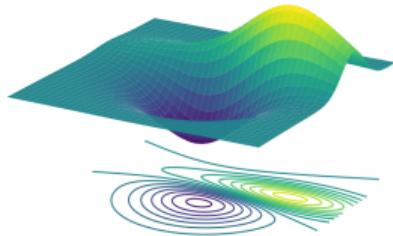
Basic Statistics



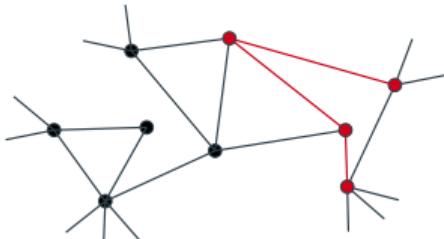
Probabilistic / Kernel Methods



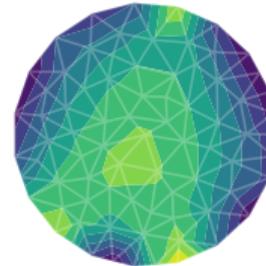
Optimization



Graphs and (Neural) Networks



Differential Equations



...and many more.



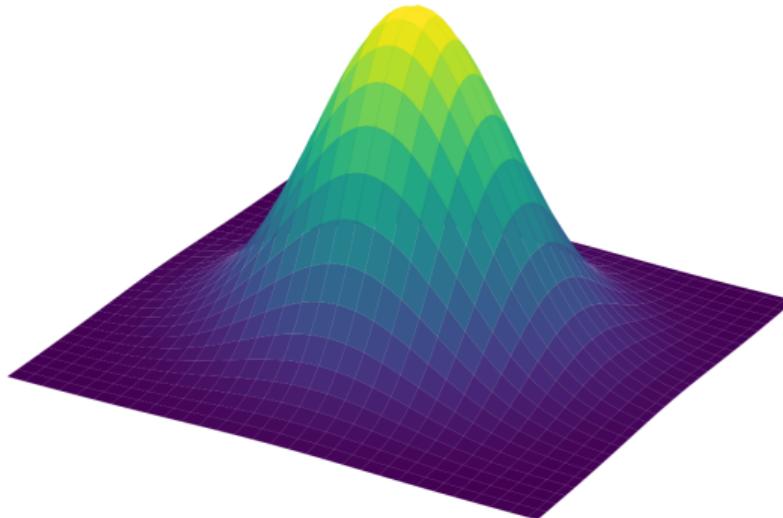
# Linear Systems are Everywhere in Scientific Computing

Example: Probability theory.

## Normal Distribution

$$x \sim \mathcal{N}(\mu, \Sigma)$$

$$p(x) = \frac{1}{\sqrt{(2\pi)^n \det(\Sigma)}} \exp\left(-\frac{1}{2}(x - \mu)^\top \Sigma^{-1} (x - \mu)\right)$$





# Linear Systems are Everywhere in Scientific Computing

Example: Probabilistic Models and Kernel Methods.

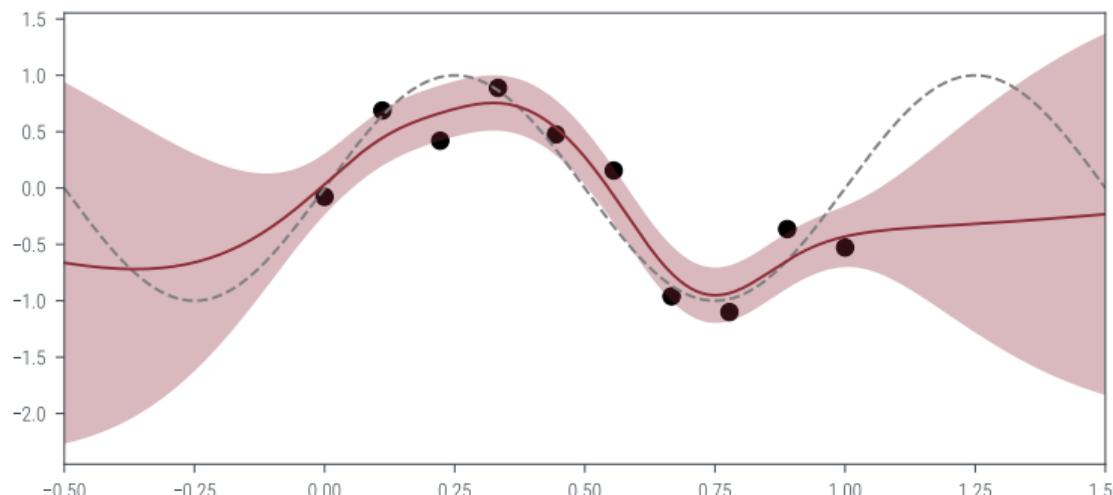
## Gaussian Processes

$$f \sim \mathcal{GP}(\mu, k)$$

$$f | X, y \sim \mathcal{GP}(\mu_{\text{post}}, k_{\text{post}})$$

$$\mu_{\text{post}}(x) = \mu(x) + k(x, X)(k(X, X) + \sigma^2 I)^{-1}(y - \mu(X))$$

$$k_{\text{post}}(x_0, x_1) = k(x_0, x_1) - k(x_0, X)(k(X, X) + \sigma^2 I)^{-1}k(X, x_1)$$



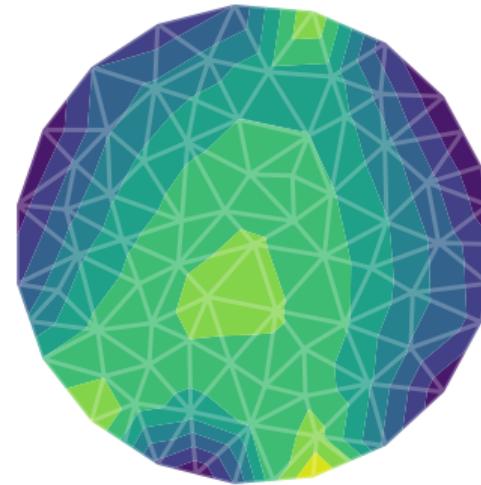
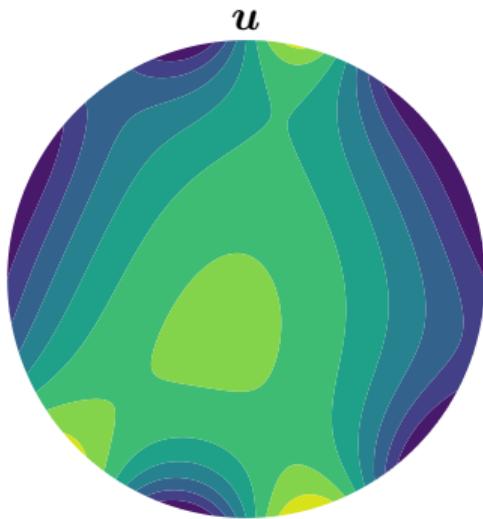


# Linear Systems are Everywhere in Scientific Computing

Example: Linear Differential Equations.

## Galerkin Method

$$\underbrace{Du = f}_{\text{linear differential equation}} \quad \Rightarrow \quad \underbrace{\hat{D}\hat{u} = \hat{f}}_{\text{finite dimensional linear system}}$$





# Linear Systems are Everywhere in Scientific Computing

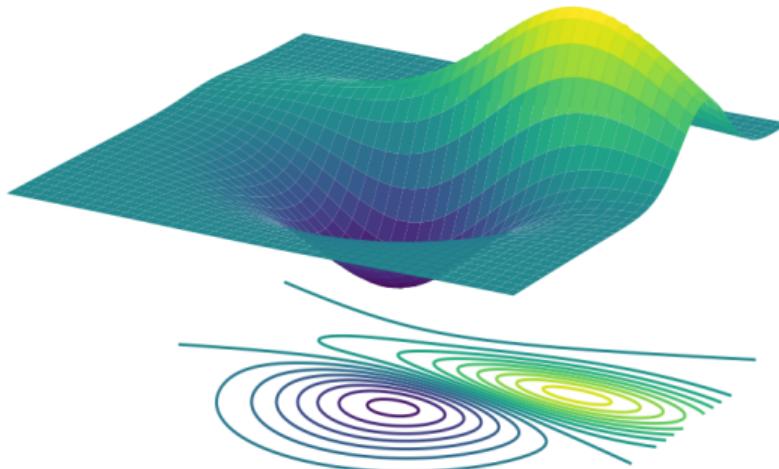
Example: Optimization.

## Iterative Optimization Methods

$$\theta_i \approx \arg \min_{\theta \in \Theta} \mathcal{L}(\theta)$$

$$\theta_i = \theta_{i-1} + \alpha_i \mathbf{M}_i \mathbf{d}_i$$

Examples: natural / conjugate / stochastic gradient descent, (Quasi-) Newton method, ...





# Linear Systems are Everywhere in Scientific Computing

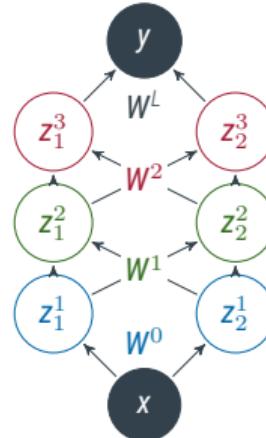
Example: Bayesian Deep Learning.

## Feedforward Neural Network

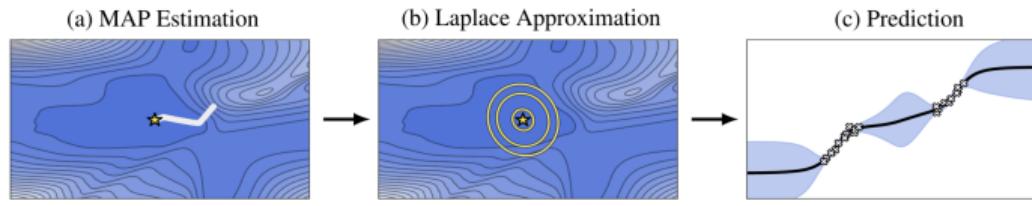
$$z^0(x, \theta) = x$$

$$z^{\ell+1}(x, \theta) = \sigma(W^\ell z^\ell + b^\ell)$$

$$y := f(x, \theta) = z^L(x, \theta)$$



Bayesian deep learning via Laplace approximation:  $p(\theta | \mathcal{D}) \approx \mathcal{N}(\theta; \theta_{\text{MAP}}, (\nabla_\theta^2 \mathcal{L}(\theta)|_{\theta_{\text{MAP}}})^{-1})$



Daxberger et al. [Dax+22]

# Probabilistic Linear Solvers

Learning the solution of a linear system.



# Probabilistic Linear Solvers for Machine Learning

Solving linear systems as probabilistic inference.

[Hen15; Coc+19b; WH20]

## Goal

Solve **large-scale** linear system  $Ax_* = b$  for  $x_* \in \mathbb{R}^n$ .

# Probabilistic Linear Solvers for Machine Learning



Solving linear systems as probabilistic inference.

[Hen15; Coc+19b; WH20]

## Goal

Solve **large-scale** linear system  $Ax_* = b$  for  $x_* \in \mathbb{R}^n$ .

## Core Insights of Probabilistic Numerics

[HOG15; Coc+19a; HOK22]

- ▶ The solution to any numerical problem is fundamentally **uncertain**.



# Probabilistic Linear Solvers for Machine Learning

Solving linear systems as probabilistic inference.

[Hen15; Coc+19b; WH20]

## Goal

Solve **large-scale** linear system  $Ax_* = b$  for  $x_* \in \mathbb{R}^n$ .

## Core Insights of Probabilistic Numerics

[HOG15; Coc+19a; HOK22]

- ▶ The solution to any numerical problem is fundamentally **uncertain**.
- ▶ Numerical algorithms are **learning agents**, which actively collect data and make predictions.



# Probabilistic Linear Solvers for Machine Learning

Solving linear systems as probabilistic inference.

[Hen15; Coc+19b; WH20]

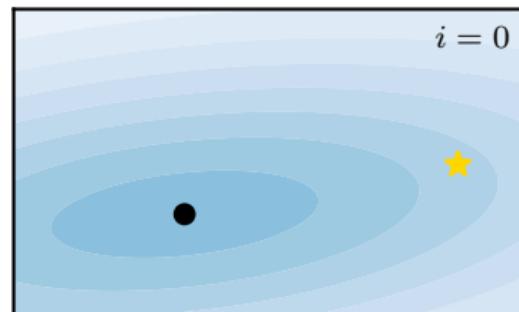
## Goal

Solve **large-scale** linear system  $\mathbf{A}\mathbf{x}_* = \mathbf{b}$  for  $\mathbf{x}_* \in \mathbb{R}^n$ .

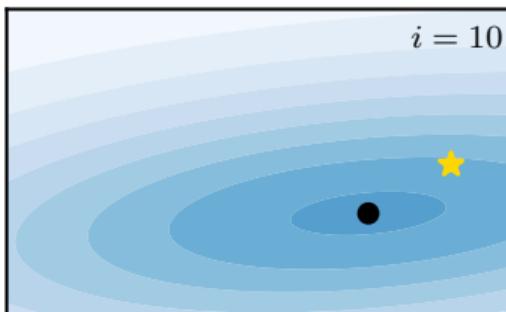
## Core Insights of Probabilistic Numerics

[HOG15; Coc+19a; HOK22]

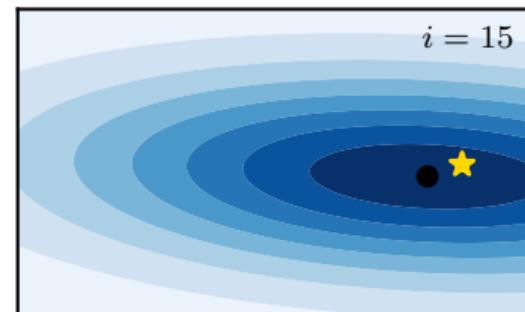
- ▶ The solution to any numerical problem is fundamentally **uncertain**.
- ▶ Numerical algorithms are **learning agents**, which actively collect data and make predictions.



Solution  $\mathbf{x}_*$



Estimate  $\mathbf{x}_i = \mathbb{E}(\mathbf{x}_*)$



Belief  $p(\mathbf{x}_*)$



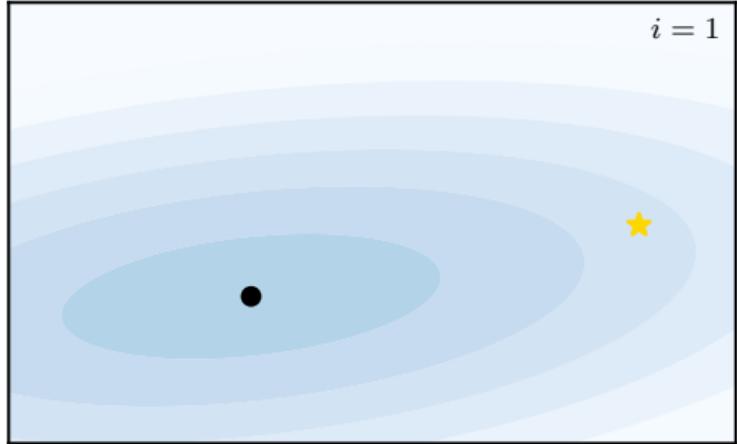
# Learning The Solution

Estimating the solution of a linear system with a probabilistic linear solver.

[Hen15; Coc+19b]

Goal: Solve  $Ax_* = b$  for  $x_*$ .

Prior:  $x_* \sim \mathcal{N}(x_0, \Sigma_0)$



- ★ Solution  $\mathbf{x}_*$
- Approximation  $\mathbf{x}_{i-1}$
- Belief  $p(\mathbf{x}_*) = \mathcal{N}(\mathbf{x}_{i-1}, \Sigma_{i-1})$



# Learning The Solution

Estimating the solution of a linear system with a probabilistic linear solver.

[Hen15; Coc+19b]

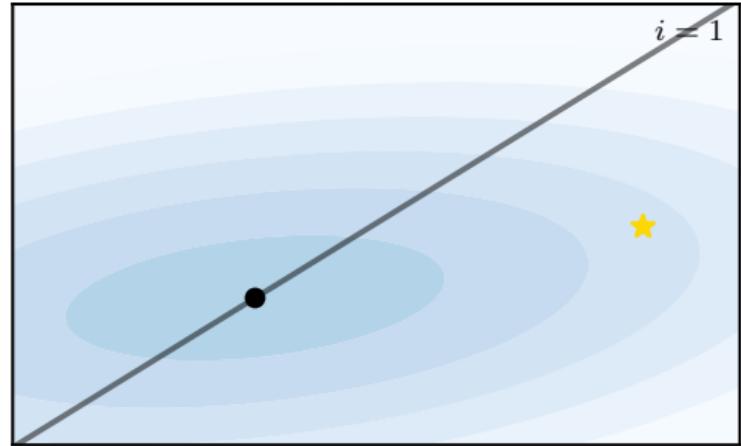
Goal: Solve  $Ax_* = b$  for  $x_*$ .

Prior:  $x_* \sim \mathcal{N}(x_0, \Sigma_0)$

Likelihood: Observe  $x_*$  via arbitrary actions  $s_i$ :

$$\alpha_i := s_i^\top A(x_* - x_{i-1}) = s_i^\top r_{i-1}$$

$$p(\alpha_i | x_*) = \lim_{\varepsilon \rightarrow 0} \mathcal{N}(\alpha_i; 0, \varepsilon)$$



- ★ Solution  $\boldsymbol{x}_*$
- Approximation  $\boldsymbol{x}_{i-1}$
- Belief  $p(\boldsymbol{x}_*) = \mathcal{N}(\boldsymbol{x}_{i-1}, \boldsymbol{\Sigma}_{i-1})$
- Action  $\boldsymbol{s}_i$



# Learning The Solution

Estimating the solution of a linear system with a probabilistic linear solver.

[Hen15; Coc+19b]

Goal: Solve  $Ax_* = b$  for  $x_*$ .

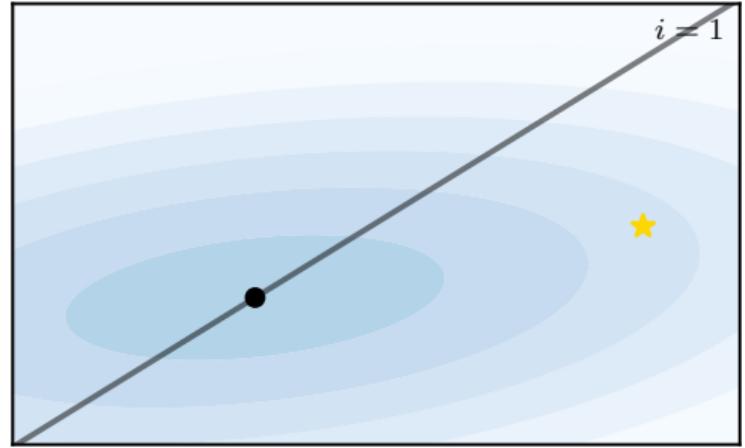
Prior:  $x_* \sim \mathcal{N}(x_0, \Sigma_0)$

Likelihood: Observe  $x_*$  via arbitrary actions  $s_i$ :

$$\alpha_i := s_i^\top A(x_* - x_{i-1}) = s_i^\top r_{i-1}$$

$$p(\alpha_i | x_*) = \lim_{\varepsilon \rightarrow 0} \mathcal{N}(\alpha_i; 0, \varepsilon)$$

Posterior: Bayes' rule gives a closed form update!



- Yellow star: Solution  $x_*$
- Black dot: Approximation  $x_{i-1}$
- Blue shaded area: Belief  $p(x_* | s_i) = \mathcal{N}(x_{i-1}, \Sigma_{i-1})$
- Black line: Action  $s_i$



# Learning The Solution

Estimating the solution of a linear system with a probabilistic linear solver.

[Hen15; Coc+19b]

Goal: Solve  $Ax_* = b$  for  $x_*$ .

Prior:  $x_* \sim \mathcal{N}(x_0, \Sigma_0)$

Likelihood: Observe  $x_*$  via arbitrary actions  $s_i$ :

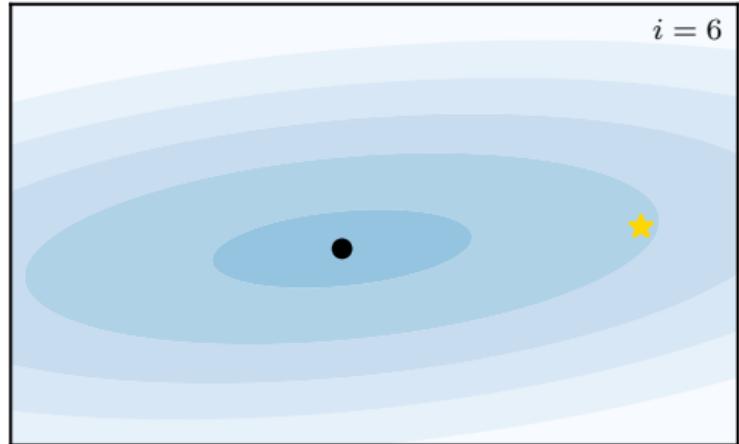
$$\alpha_i := s_i^\top A(x_* - x_{i-1}) = s_i^\top r_{i-1}$$

$$p(\alpha_i | x_*) = \lim_{\varepsilon \rightarrow 0} \mathcal{N}(\alpha_i; 0, \varepsilon)$$

Posterior:  $x_* | \alpha_1, \dots, \alpha_i \sim \mathcal{N}(x_i, \Sigma_i)$

$$x_i = x_0 + \Sigma_0 A S_i (S_i^\top A \Sigma_0 A S_i)^{-1} S_i^\top (b - Ax_0)$$

$$\Sigma_i = \Sigma_0 - \Sigma_0 A S_i (S_i^\top A \Sigma_0 A S_i)^{-1} S_i^\top A \Sigma_0$$



★ Solution  $\boldsymbol{x}_*$

● Approximation  $\boldsymbol{x}_{i-1}$

■ Belief  $p(\boldsymbol{x}_*) = \mathcal{N}(\boldsymbol{x}_{i-1}, \boldsymbol{\Sigma}_{i-1})$



# Learning The Solution

Estimating the solution of a linear system with a probabilistic linear solver.

[Hen15; Coc+19b]

Goal: Solve  $Ax_* = b$  for  $x_*$ .

Prior:  $x_* \sim \mathcal{N}(x_0, \Sigma_0)$

Likelihood: Observe  $x_*$  via arbitrary actions  $s_i$ :

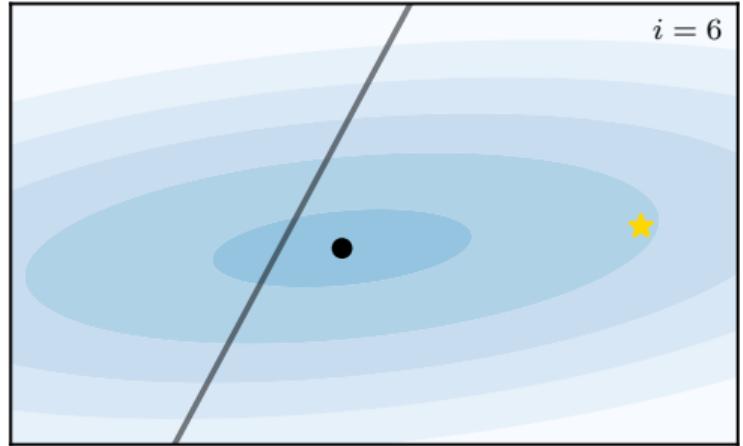
$$\alpha_i := s_i^\top A(x_* - x_{i-1}) = s_i^\top r_{i-1}$$

$$p(\alpha_i | x_*) = \lim_{\varepsilon \rightarrow 0} \mathcal{N}(\alpha_i; 0, \varepsilon)$$

Posterior:  $x_* | \alpha_1, \dots, \alpha_i \sim \mathcal{N}(x_i, \Sigma_i)$

$$x_i = x_0 + \Sigma_0 A S_i (S_i^\top A \Sigma_0 A S_i)^{-1} S_i^\top (b - Ax_0)$$

$$\Sigma_i = \Sigma_0 - \Sigma_0 A S_i (S_i^\top A \Sigma_0 A S_i)^{-1} S_i^\top A \Sigma_0$$



Solution  $\boldsymbol{x}_*$



Approximation  $\boldsymbol{x}_{i-1}$



Belief  $p(\boldsymbol{x}_*) = \mathcal{N}(\boldsymbol{x}_{i-1}, \boldsymbol{\Sigma}_{i-1})$



Action  $\boldsymbol{s}_i$



# Learning The Solution

Estimating the solution of a linear system with a probabilistic linear solver.

[Hen15; Coc+19b]

Goal: Solve  $Ax_* = b$  for  $x_*$ .

Prior:  $x_* \sim \mathcal{N}(x_0, \Sigma_0)$

Likelihood: Observe  $x_*$  via arbitrary actions  $s_i$ :

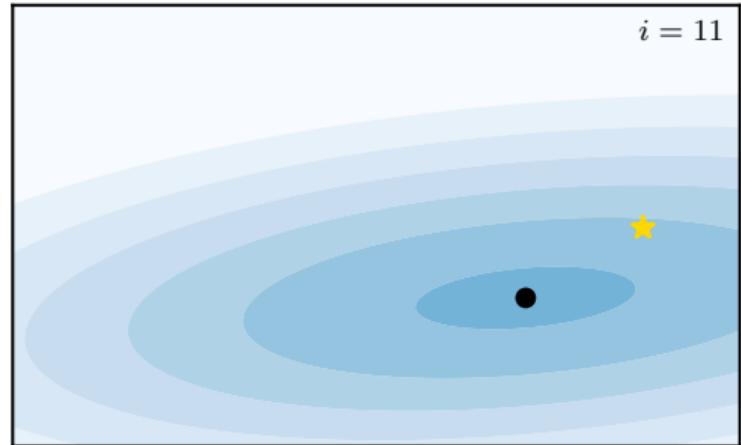
$$\alpha_i := s_i^\top A(x_* - x_{i-1}) = s_i^\top r_{i-1}$$

$$p(\alpha_i | x_*) = \lim_{\varepsilon \rightarrow 0} \mathcal{N}(\alpha_i; 0, \varepsilon)$$

Posterior:  $x_* | \alpha_1, \dots, \alpha_i \sim \mathcal{N}(x_i, \Sigma_i)$

$$x_i = x_0 + \Sigma_0 A S_i (S_i^\top A \Sigma_0 A S_i)^{-1} S_i^\top (b - Ax_0)$$

$$\Sigma_i = \Sigma_0 - \Sigma_0 A S_i (S_i^\top A \Sigma_0 A S_i)^{-1} S_i^\top A \Sigma_0$$



★ Solution  $\boldsymbol{x}_*$

● Approximation  $\boldsymbol{x}_{i-1}$

■ Belief  $p(\boldsymbol{x}_*) = \mathcal{N}(\boldsymbol{x}_{i-1}, \boldsymbol{\Sigma}_{i-1})$



# Learning The Solution

Estimating the solution of a linear system with a probabilistic linear solver.

[Hen15; Coc+19b]

Goal: Solve  $Ax_* = b$  for  $x_*$ .

Prior:  $x_* \sim \mathcal{N}(x_0, \Sigma_0)$

Likelihood: Observe  $x_*$  via arbitrary actions  $s_i$ :

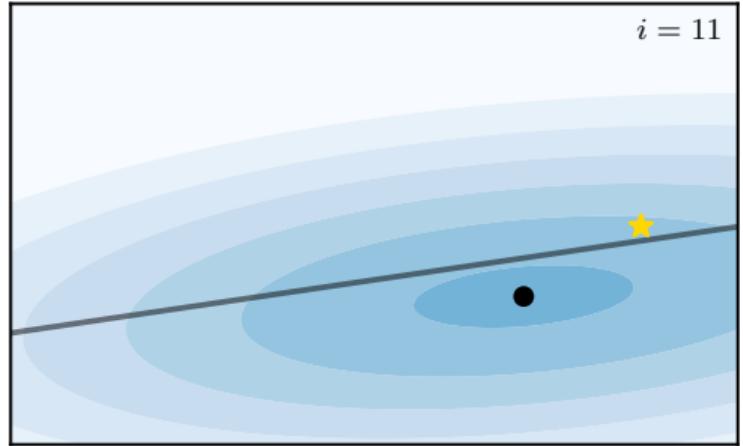
$$\alpha_i := s_i^\top A(x_* - x_{i-1}) = s_i^\top r_{i-1}$$

$$p(\alpha_i | x_*) = \lim_{\varepsilon \rightarrow 0} \mathcal{N}(\alpha_i; 0, \varepsilon)$$

Posterior:  $x_* | \alpha_1, \dots, \alpha_i \sim \mathcal{N}(x_i, \Sigma_i)$

$$x_i = x_0 + \Sigma_0 A S_i (S_i^\top A \Sigma_0 A S_i)^{-1} S_i^\top (b - Ax_0)$$

$$\Sigma_i = \Sigma_0 - \Sigma_0 A S_i (S_i^\top A \Sigma_0 A S_i)^{-1} S_i^\top A \Sigma_0$$



Solution  $\boldsymbol{x}_*$



Approximation  $\boldsymbol{x}_{i-1}$



Belief  $p(\boldsymbol{x}_*) = \mathcal{N}(\boldsymbol{x}_{i-1}, \boldsymbol{\Sigma}_{i-1})$



Action  $\boldsymbol{s}_i$



# Learning The Solution

Estimating the solution of a linear system with a probabilistic linear solver.

[Hen15; Coc+19b]

Goal: Solve  $Ax_* = b$  for  $x_*$ .

Prior:  $x_* \sim \mathcal{N}(x_0, \Sigma_0)$

Likelihood: Observe  $x_*$  via arbitrary actions  $s_i$ :

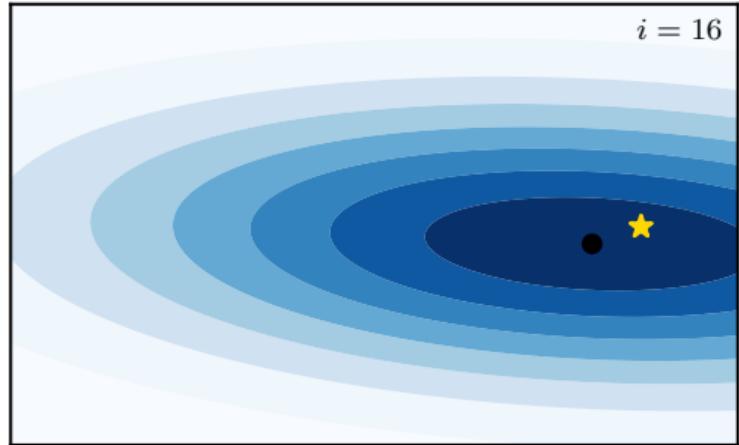
$$\alpha_i := s_i^\top A(x_* - x_{i-1}) = s_i^\top r_{i-1}$$

$$p(\alpha_i | x_*) = \lim_{\varepsilon \rightarrow 0} \mathcal{N}(\alpha_i; 0, \varepsilon)$$

Posterior:  $x_* | \alpha_1, \dots, \alpha_i \sim \mathcal{N}(x_i, \Sigma_i)$

$$x_i = x_0 + \Sigma_0 A S_i (S_i^\top A \Sigma_0 A S_i)^{-1} S_i^\top (b - Ax_0)$$

$$\Sigma_i = \Sigma_0 - \Sigma_0 A S_i (S_i^\top A \Sigma_0 A S_i)^{-1} S_i^\top A \Sigma_0$$



- ★ Solution  $\mathbf{x}_*$
- Approximation  $\mathbf{x}_{i-1}$
- Belief  $p(\mathbf{x}_*) = \mathcal{N}(\mathbf{x}_{i-1}, \Sigma_{i-1})$



# Learning The Solution

Estimating the solution of a linear system with a probabilistic linear solver.

[Hen15; Coc+19b]

Goal: Solve  $Ax_* = b$  for  $x_*$ .

Prior:  $x_* \sim \mathcal{N}(x_0, \Sigma_0)$

Likelihood: Observe  $x_*$  via arbitrary actions  $s_i$ :

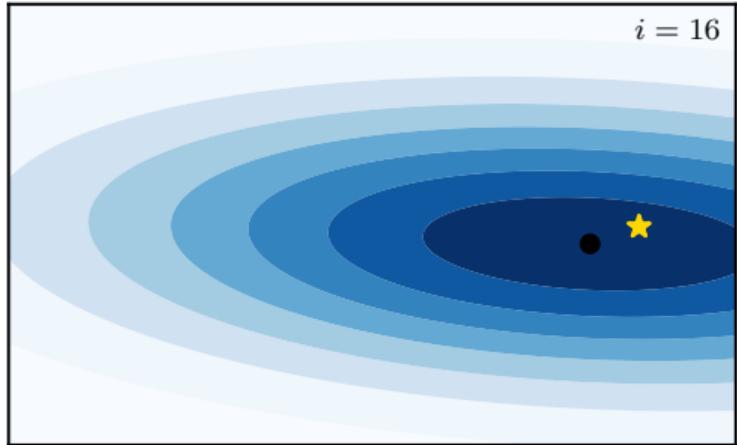
$$\alpha_i := s_i^\top A(x_* - x_{i-1}) = s_i^\top r_{i-1}$$

$$p(\alpha_i | x_*) = \lim_{\varepsilon \rightarrow 0} \mathcal{N}(\alpha_i; 0, \varepsilon)$$

Posterior:  $x_* | \alpha_1, \dots, \alpha_i \sim \mathcal{N}(x_i, \Sigma_i)$

$$x_i = x_0 + \Sigma_0 A S_i (S_i^\top A \Sigma_0 A S_i)^{-1} S_i^\top (b - Ax_0)$$

$$\Sigma_i = \Sigma_0 - \Sigma_0 A S_i (S_i^\top A \Sigma_0 A S_i)^{-1} S_i^\top A \Sigma_0$$



- ★ Solution  $\mathbf{x}_*$
- Approximation  $\mathbf{x}_{i-1}$
- Belief  $p(\mathbf{x}_*) = \mathcal{N}(\mathbf{x}_{i-1}, \Sigma_{i-1})$

How do we choose the linear solver actions  $S$  and the prior  $\mathcal{N}(x_0, \Sigma_0)$ ?



# Policy Choice

How do we choose the actions?

**Observation:** Actions “weigh” entries in the residual:  $\alpha_i := \mathbf{s}_i^\top \mathbf{r}_{i-1} = \mathbf{s}_i^\top \mathbf{A}(\mathbf{x}_* - \mathbf{x}_{i-1})$



# Policy Choice

How do we choose the actions?

**Observation:** Actions “weigh” entries in the residual:  $\alpha_i := \mathbf{s}_i^\top \mathbf{r}_{i-1} = \mathbf{s}_i^\top \mathbf{A}(\mathbf{x}_* - \mathbf{x}_{i-1})$

**Idea:** Focus computation where residual is large:  $\mathbf{s}_i = \mathbf{r}_{i-1} \Rightarrow \alpha_i = \|\mathbf{r}_{i-1}\|_2^2$   
 $\Rightarrow$  **BayesCG** [Coc+19b]



# Interlude: Method of Conjugate Gradients

Efficiently solving linear systems with positive definite system matrix via matrix-vector multiplies.

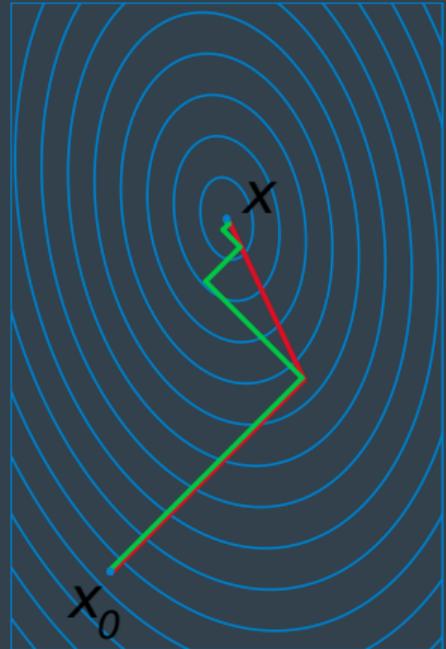
**Goal:** Approximately solve linear system  $Ax_* = b$ .

**Idea:** Rephrase as quadratic optimization problem and optimize. Let

$$f(x) = \frac{1}{2}x^\top Ax - b^\top x$$

then  $\nabla f(x_*) = \mathbf{0} \iff Ax_* = b \iff r(x_*) := b - Ax_* = \mathbf{0}$ .

**Question:** How should we optimize?



Oleg Alexandrov, commons.wikimedia.org/w/index.php?curid=2267598



# Interlude: Method of Conjugate Gradients

Efficiently solving linear systems with positive definite system matrix via matrix-vector multiplies.

**Goal:** Approximately solve linear system  $\mathbf{A}\mathbf{x}_* = \mathbf{b}$ .

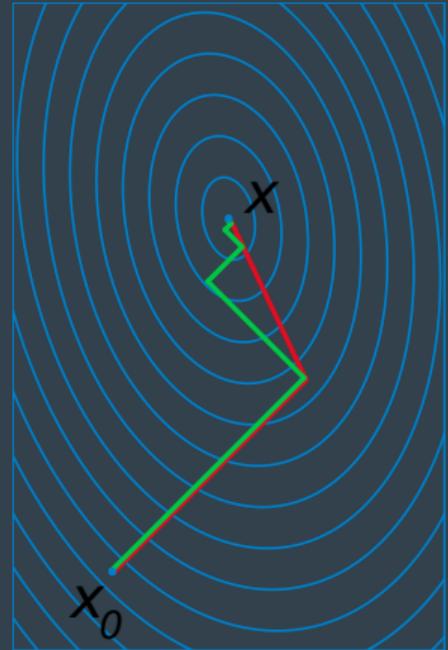
**Idea:** Rephrase as quadratic optimization problem and optimize. Let

$$f(\mathbf{x}) = \frac{1}{2} \mathbf{x}^\top \mathbf{A} \mathbf{x} - \mathbf{b}^\top \mathbf{x}$$

then  $\nabla f(\mathbf{x}_*) = \mathbf{0} \iff \mathbf{A}\mathbf{x}_* = \mathbf{b} \iff \mathbf{r}(\mathbf{x}_*) := \mathbf{b} - \mathbf{A}\mathbf{x}_* = \mathbf{0}$ .

**Question:** How should we optimize?

1. **Gradient descent:** Follow  $\mathbf{d}_i = \mathbf{r}(\mathbf{x}_i) = -\nabla f(\mathbf{x}_i)$  s.t.  $\langle \mathbf{d}_i, \mathbf{d}_j \rangle = 0$ .



Oleg Alexandrov, commons.wikimedia.org/w/index.php?curid=2267598



# Interlude: Method of Conjugate Gradients

Efficiently solving linear systems with positive definite system matrix via matrix-vector multiplies.

**Goal:** Approximately solve linear system  $Ax_* = b$ .

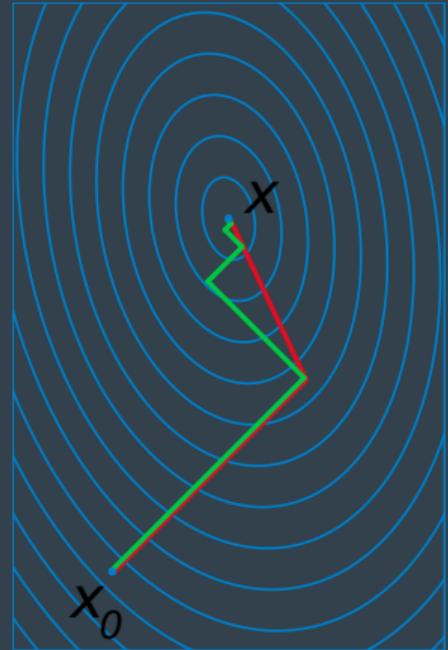
**Idea:** Rephrase as quadratic optimization problem and optimize. Let

$$f(x) = \frac{1}{2}x^\top Ax - b^\top x$$

then  $\nabla f(x_*) = \mathbf{0} \iff Ax_* = b \iff r(x_*) := b - Ax_* = \mathbf{0}$ .

**Question:** How should we optimize?

- 1 **Gradient descent:** Follow  $d_i = r(x_i) = -\nabla f(x_i)$  s.t.  $\langle d_i, d_j \rangle = 0$ .
- 2 Conjugate direction method: Follow  $d_i$  s. t.  $\langle d_i^\top d_j \rangle_A = d_i^\top A d_j = 0$  for  $i \neq j$ .  
 $\Rightarrow$  convergence in at most  $n$  steps.



Oleg Alexandrov, commons.wikimedia.org/w/index.php?curid=2267598



# Interlude: Method of Conjugate Gradients

Efficiently solving linear systems with positive definite system matrix via matrix-vector multiplies.

**Goal:** Approximately solve linear system  $Ax_* = b$ .

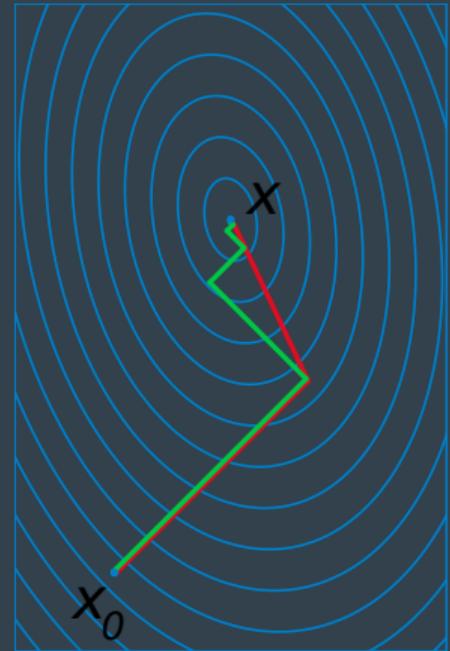
**Idea:** Rephrase as quadratic optimization problem and optimize. Let

$$f(x) = \frac{1}{2}x^\top Ax - b^\top x$$

then  $\nabla f(x_*) = \mathbf{0} \iff Ax_* = b \iff r(x_*) := b - Ax_* = \mathbf{0}$ .

**Question:** How should we optimize?

- 1 **Gradient descent:** Follow  $d_i = r(x_i) = -\nabla f(x_i)$  s.t.  $\langle d_i, d_j \rangle = 0$ .
- 2 Conjugate direction method: Follow  $d_i$  s. t.  $\langle d_i^\top d_j \rangle_A = d_i^\top A d_j = 0$  for  $i \neq j$ .  
 $\Rightarrow$  convergence in at most  $n$  steps.
- 3 **Conjugate gradient method:** First step  $d_0 = r(x_0)$ .



Oleg Alexandrov, commons.wikimedia.org/w/index.php?curid=2267598



# Policy Choice

How do we choose the actions?

**Observation:** Actions  $s_i$  "weigh" entries in the residual:  $\alpha_i := s_i^\top r_{i-1} = s_i^\top A(x_* - x_{i-1})$

**Idea:** Focus computation where residual is large:  $s_i = r_{i-1} \Rightarrow \alpha_i = \|r_{i-1}\|_2^2$   
 $\Rightarrow$  **BayesCG** [Coc+19b]

Theorem (Equivalence to Conjugate Gradient Method [Coc+19b; Wen+22])

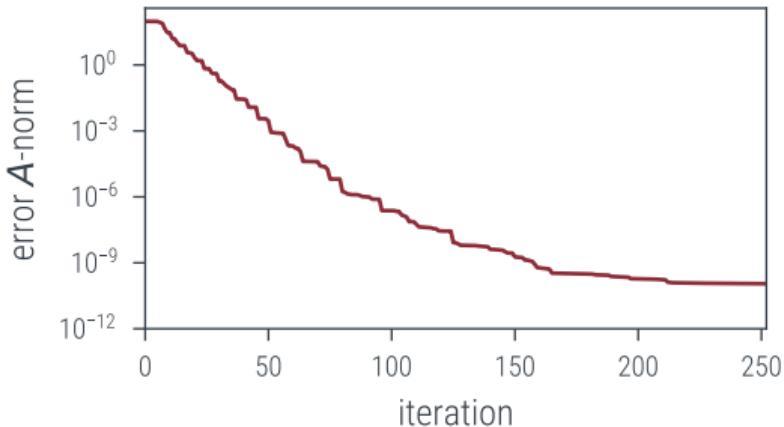
If  $x_0 = \mathbf{0}$ ,  $\Sigma_0 = A^{-1}$  and the actions are either conjugate gradients  $s_i = d_i^{CG}$  or gradients  $s_i = r_{i-1}$ , then the posterior mean  $x_i = x_i^{CG}$  of BayesCG is equivalent to the approximation returned by CG.



# Convergence Behavior of the Conjugate Gradient Method

The spectrum of the matrix determines the convergence speed.

$$n = 10^3 \quad \kappa(A) \approx 7 \cdot 10^5$$



Theorem (Convergence Rate of CG[TB97])

$$\|x - x_i\|_A \leq 2 \left( \frac{\sqrt{\kappa(A)} - 1}{\sqrt{\kappa(A)} + 1} \right)^i \|x - x_0\|_A$$

CG converges fast for a small condition number.

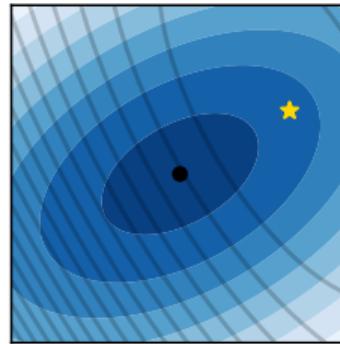
# Prior Choice

Comparing different choices of prior for BayesCG.

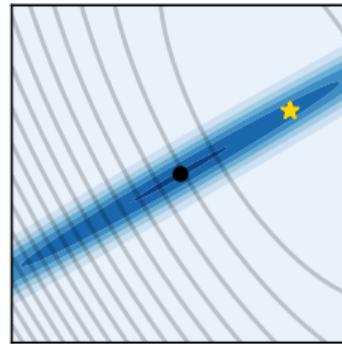
[Coc+19b]

Prior

$$x_* \sim \mathcal{N}(x_0, \Sigma_0)$$



⇒



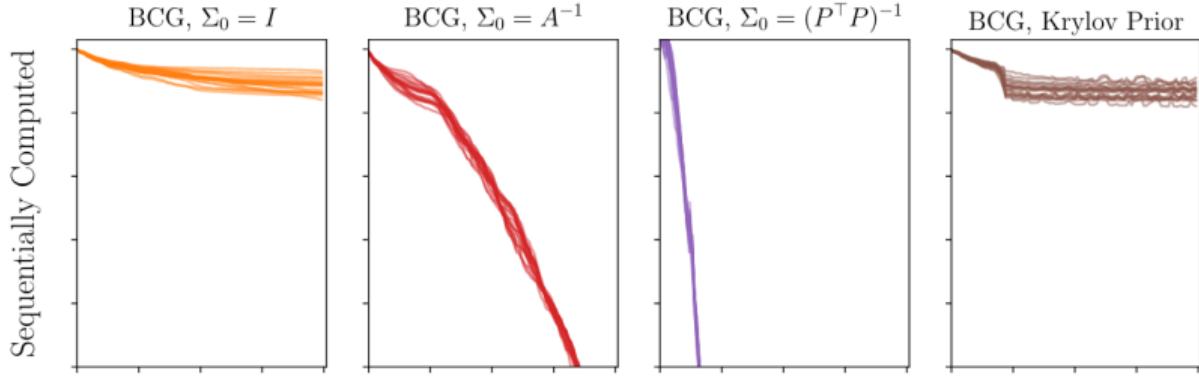
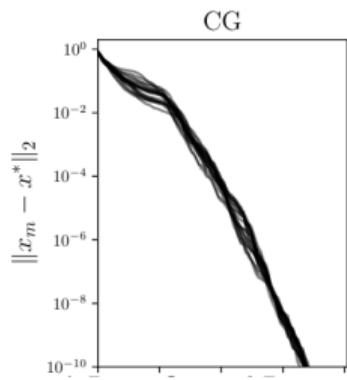
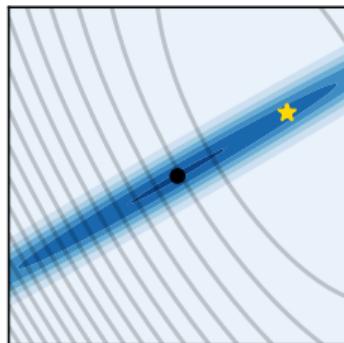
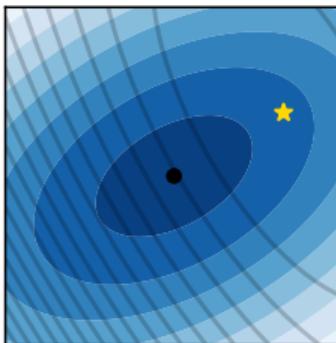
# Prior Choice

Comparing different choices of prior for BayesCG.

[Coc+19b]

Prior

$$x_* \sim \mathcal{N}(x_0, \Sigma_0)$$





# Algorithm: Probabilistic Linear Solver

Sequential formulation.

---

		Time	Space
1	procedure PROBABILISTICLINEARSOLVER( $A, b, x_0 = \mathbf{0}, \Sigma_0$ )		
2	while not STOPPINGCRITERION() do		
3	$s_i \leftarrow \text{POLICY}()$	Select action via policy.	
4	$r_{i-1} \leftarrow b - Ax_{i-1}$	Residual.	$\mathcal{O}(n^2)$
5	$\alpha_i \leftarrow s_i^\top r_{i-1}$	Observation.	$\mathcal{O}(n)$
6	$z_i \leftarrow As_i$		$\mathcal{O}(n^2)$
7	$d_i \leftarrow \Sigma_{i-1}As_i = \Sigma_{i-1}z_i$	Search direction.	$\mathcal{O}(n^2)$
8	$\eta_i \leftarrow s_i^\top A \Sigma_{i-1} As_i = z_i^\top d_i$		$\mathcal{O}(n)$
9	$C_i \leftarrow C_{i-1} + \frac{1}{\eta_i} d_i d_i^\top$		$\mathcal{O}(n)$
10	$x_i \leftarrow x_{i-1} + \frac{\alpha_i}{\eta_i} d_i$	Solution estimate.	$\mathcal{O}(n)$
11	$\Sigma_i \leftarrow \Sigma_0 - C_i$	Uncertainty.	
12	return $\mathcal{N}(x_i, \Sigma_i)$		

---

# Application: Gaussian Processes

Scaling Gaussian processes via probabilistic linear solvers.



# Gaussian Process Regression

Learning an unknown function from data.

**Goal:** Supervised learning from  $n$  data points  $(X, y)$

**Prior:** Gaussian process  $f \sim \mathcal{GP}(\mu, k)$

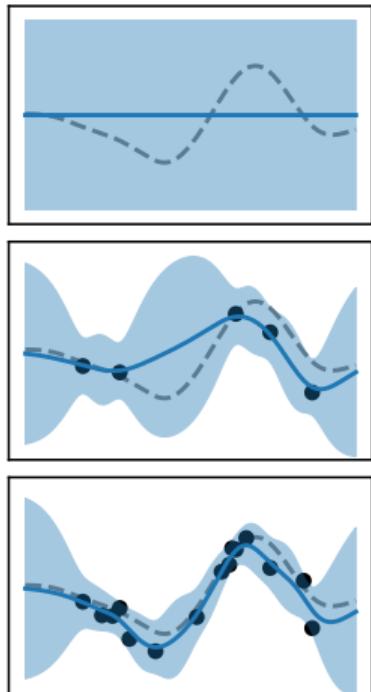
**Likelihood:** Observations  $y = f(X) + \varepsilon \sim \mathcal{N}(f(X), \sigma^2 I)$

**Posterior:**  $f | X, y \sim \mathcal{GP}(\mu_*, k_*)$  with

$$\mu_*(\cdot) = \mu(\cdot) + K(\cdot, X)\hat{K}^{-1}(y - \mu(X))$$

$$K_*(\cdot, \cdot) = K(\cdot, \cdot) - K(\cdot, X)\hat{K}^{-1}K(X, \cdot)$$

where  $\hat{K} = K + \sigma^2 I \in \mathbb{R}^{n \times n}$ .





# Gaussian Process Regression

Learning an unknown function from data.

Goal: Supervised learning from  $n$  data points  $(X, y)$

Prior: Gaussian process  $f \sim \mathcal{GP}(\mu, k)$

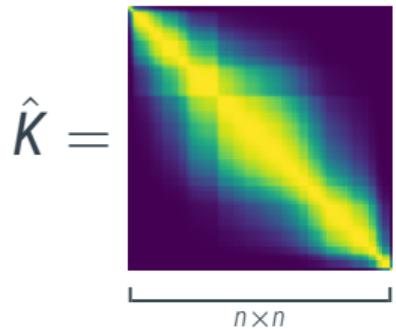
Likelihood: Observations  $y = f(X) + \varepsilon \sim \mathcal{N}(f(X), \sigma^2 I)$

Posterior:  $f | X, y \sim \mathcal{GP}(\mu_*, k_*)$  with

$$\mu_*(\cdot) = \mu(\cdot) + K(\cdot, X) \hat{K}^{-1} (y - \mu(X))$$

$$K_*(\cdot, \cdot) = K(\cdot, \cdot) - K(\cdot, X) \hat{K}^{-1} K(X, \cdot)$$

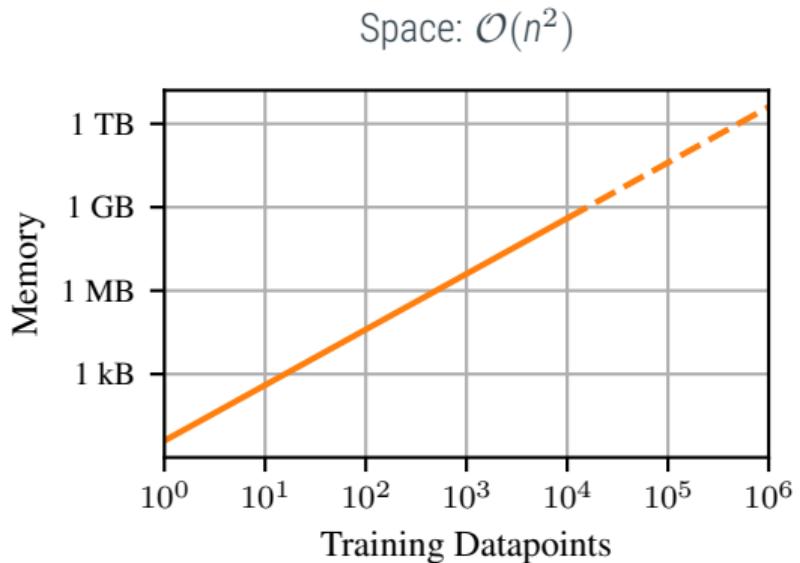
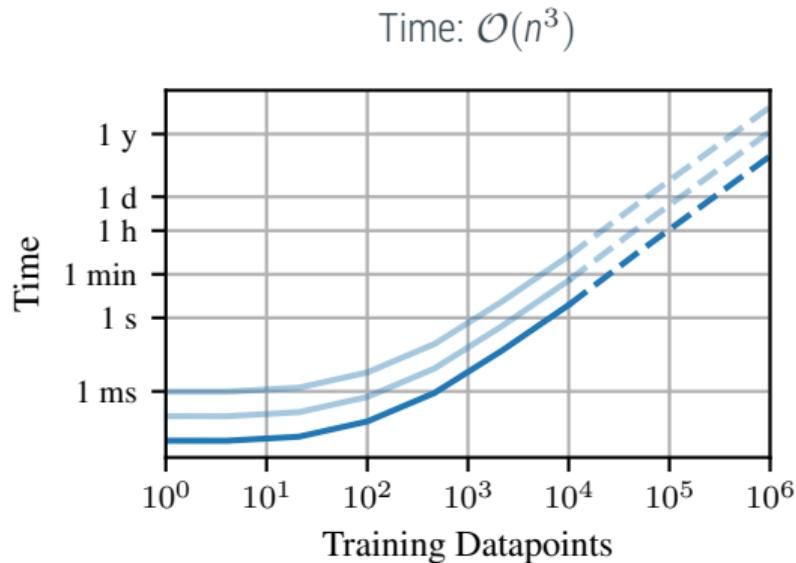
where  $\hat{K} = K + \sigma^2 I \in \mathbb{R}^{n \times n}$ .





# Computational Cost of Gaussian Processes

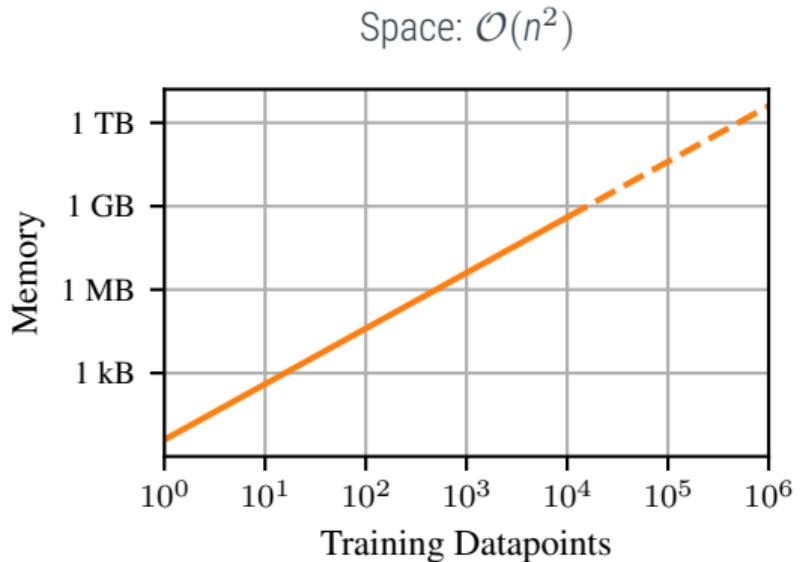
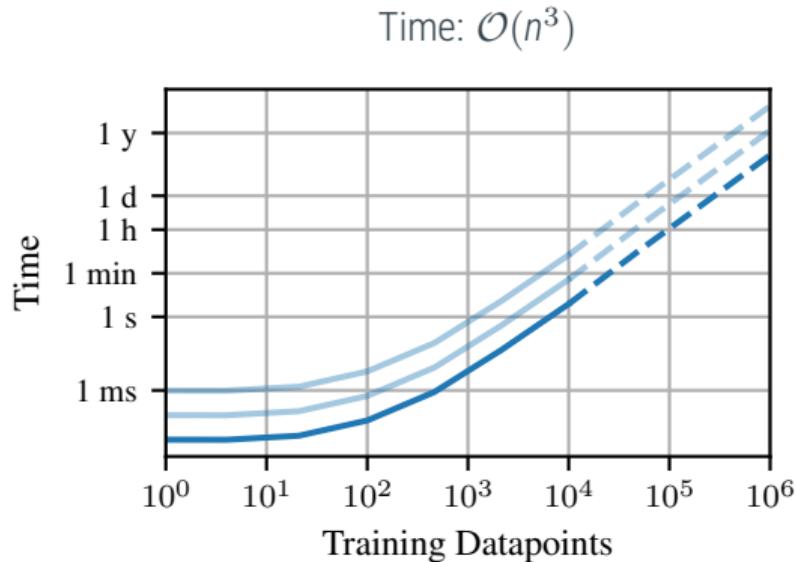
Gaussian processes scale prohibitively with the size  $n$  of the dataset.





# Computational Cost of Gaussian Processes

Gaussian processes scale prohibitively with the size  $n$  of the dataset.



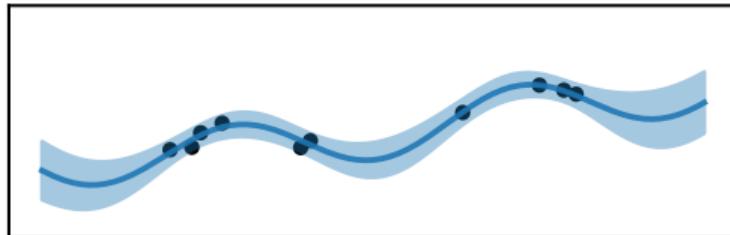
We need to **approximate** the posterior.



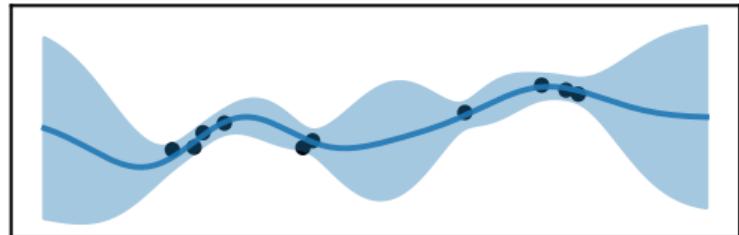
# Approximate Gaussian Process Inference

Impact of approximations on uncertainty quantification and decision-making.

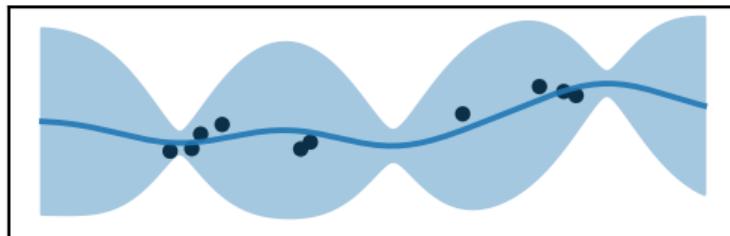
RFFGP



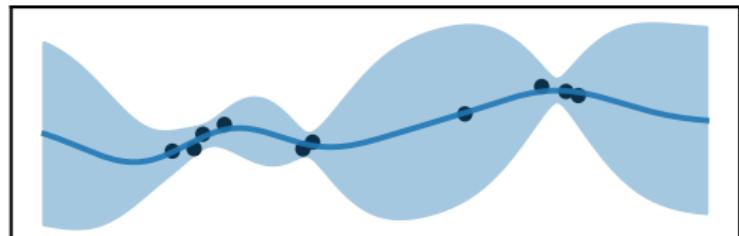
CGGP



SVGP-fixed



SVGP-opt



● Data

— Approx. Posterior Mean

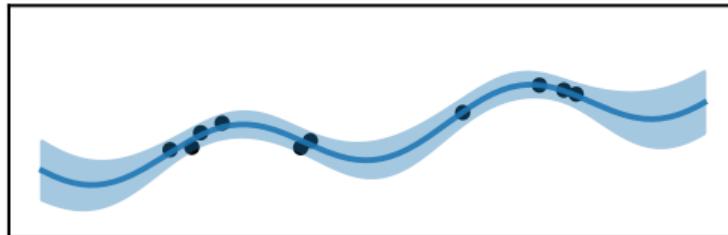
■ Approx. Posterior Uncertainty



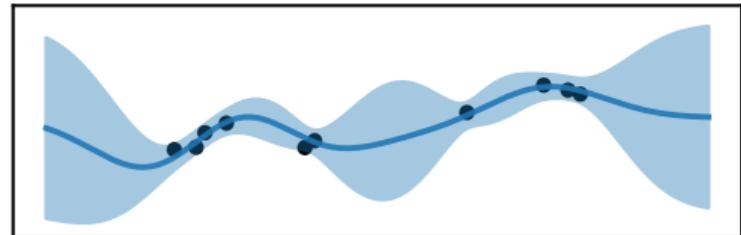
# Approximate Gaussian Process Inference

Impact of approximations on uncertainty quantification and decision-making.

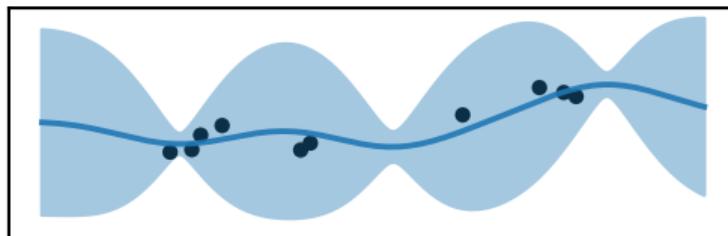
RFFGP



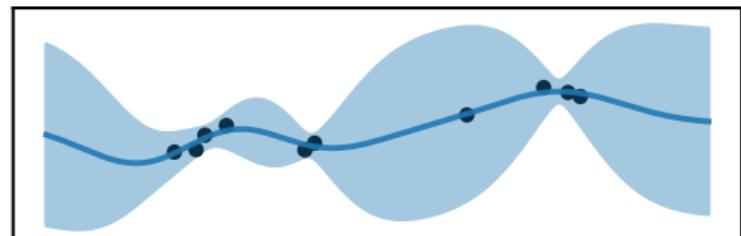
CGGP



SVGP-fixed



SVGP-opt



● Data

— Approx. Posterior Mean

■ Approx. Posterior Uncertainty

Approximations introduce **error**, which **impacts downstream decisions**.



# Fundamental Questions

## Question 1:

How can we perform Gaussian process inference at scale?



# Fundamental Questions

## Question 1:

How can we perform Gaussian process inference at scale?

## Question 2:

How can we quantify the inevitable approximation error?

# Q1: Gaussian Process Inference at Scale?

Efficiently approximating the posterior of a Gaussian process.

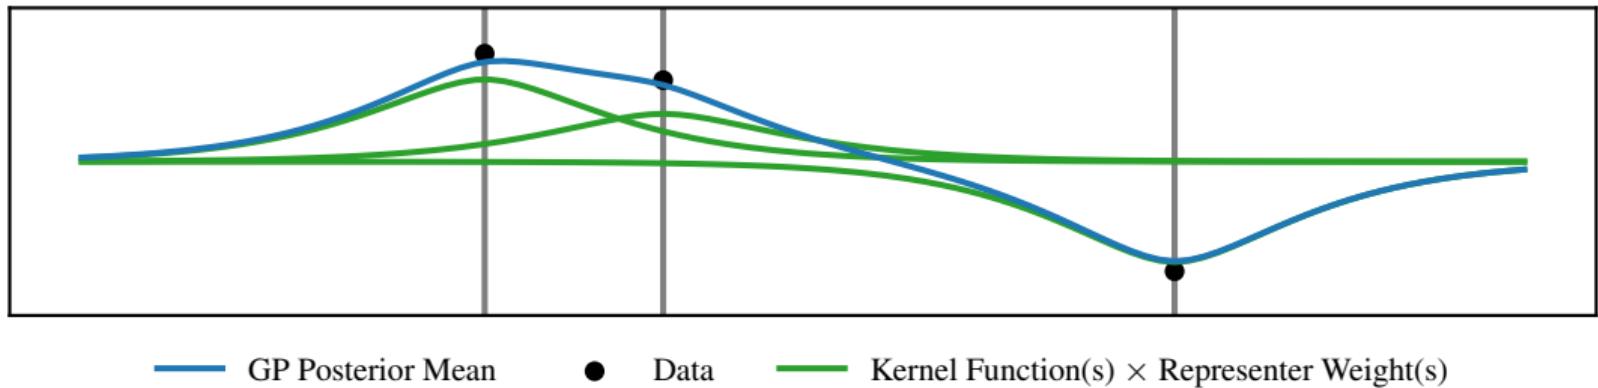


# Representer Weights

The posterior mean is a linear combination of kernel functions centered at data points.

$$f | X, y \sim \mathcal{GP}(\mu_*, k_*)$$

$$\mu_*(\cdot) = \mu(\cdot) + k(\cdot, X) \underbrace{\hat{K}^{-1}(y - \mu(X))}_{\text{representer weights } v_*} = \mu(\cdot) + \sum_{j=1}^n k(\cdot, x_j)(v_*)_j$$



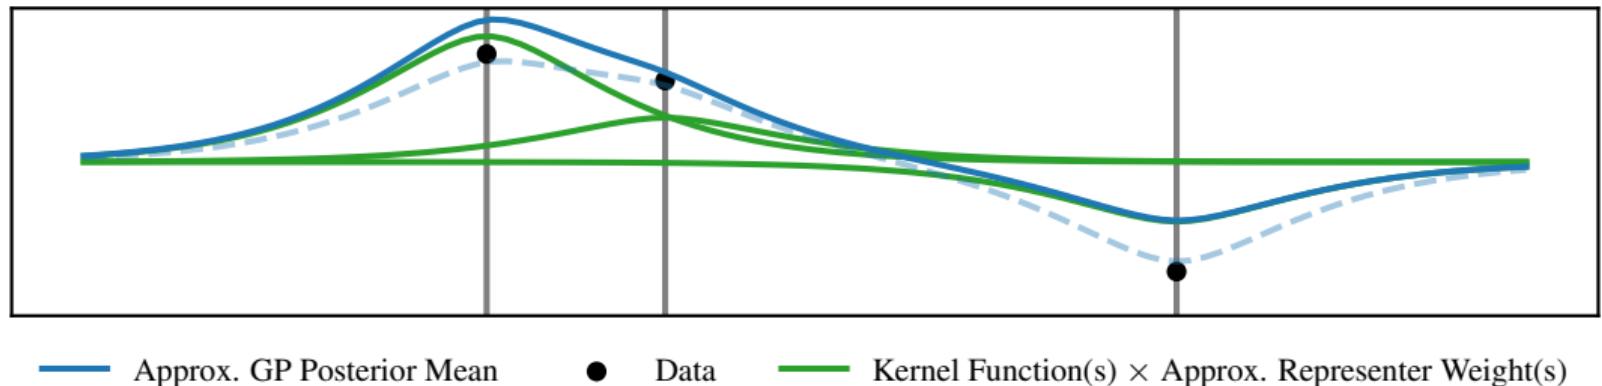


# Approximating Representer Weights

Iterative linear solvers can be used to approximate the representer weights.

$$\mu_*(\cdot) = \mu(\cdot) + k(\cdot, X) \underbrace{\hat{K}^{-1}(y - \mu(X))}_{\text{representer weights } v_*} \approx \mu(\cdot) + k(\cdot, X)v_i$$

Known: Can use iterative linear solvers (e.g. CG) to approximate representer weights  $v_* \approx v_i$ .



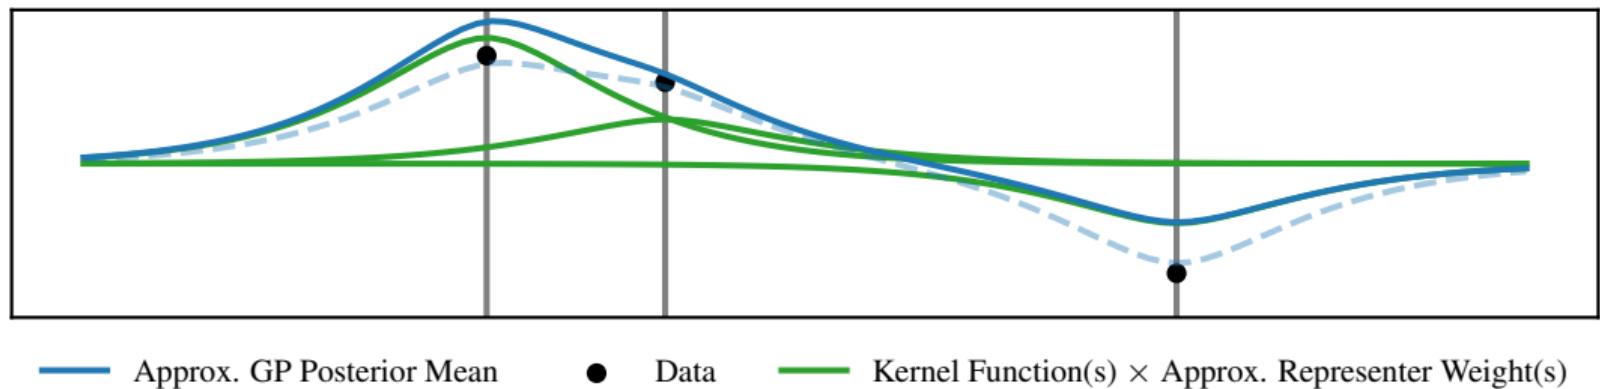


# Approximating Representer Weights

Iterative linear solvers can be used to approximate the representer weights.

$$\mu_*(\cdot) = \mu(\cdot) + k(\cdot, X) \underbrace{\hat{K}^{-1}(y - \mu(X))}_{\text{representer weights } v_*} \approx \mu(\cdot) + k(\cdot, X)v_i$$

Known: Can use iterative linear solvers (e.g. CG) to approximate representer weights  $v_* \approx v_i$ .



— Approx. GP Posterior Mean

● Data

— Kernel Function(s)  $\times$  Approx. Representer Weight(s)

Benefit: Time complexity  $\mathcal{O}(n^2)$  and space complexity  $\mathcal{O}(nd)$ .

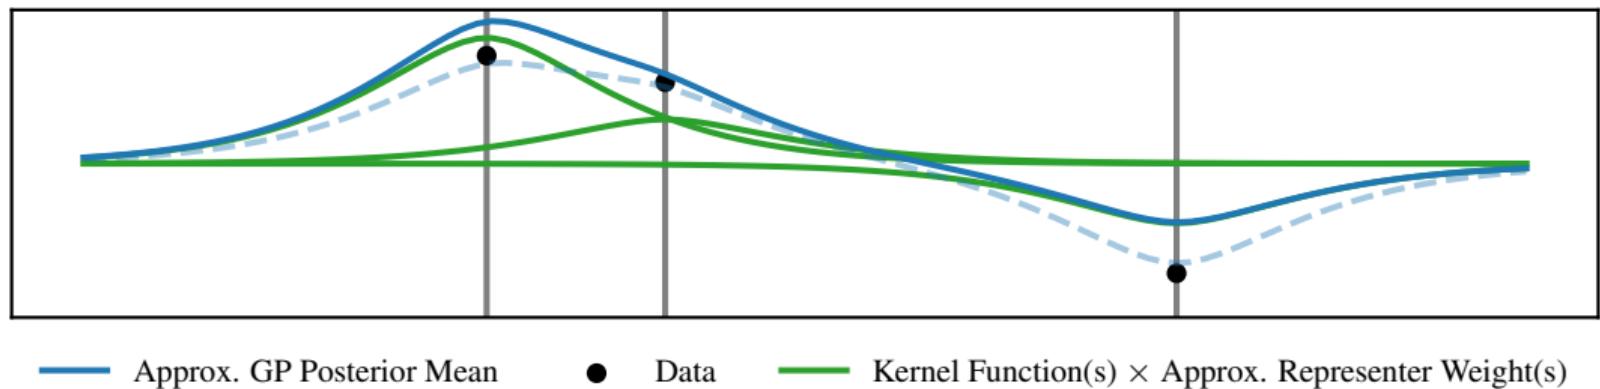


# Approximating Representer Weights

Iterative linear solvers can be used to approximate the representer weights.

$$\mu_*(\cdot) = \mu(\cdot) + k(\cdot, X) \underbrace{\hat{K}^{-1}(y - \mu(X))}_{\text{representer weights } v_*} \approx \mu(\cdot) + k(\cdot, X)v_i$$

Known: Can use iterative linear solvers (e.g. CG) to approximate representer weights  $v_* \approx v_i$ .



— Approx. GP Posterior Mean

● Data

— Kernel Function(s)  $\times$  Approx. Representer Weight(s)

Problem: Approximation error of the linear solve.

## Q2: Can We Quantify Approximation Error?

Probabilistic error quantification at prediction time using probabilistic linear solvers.



# Linear Solver Prior for GP Inference

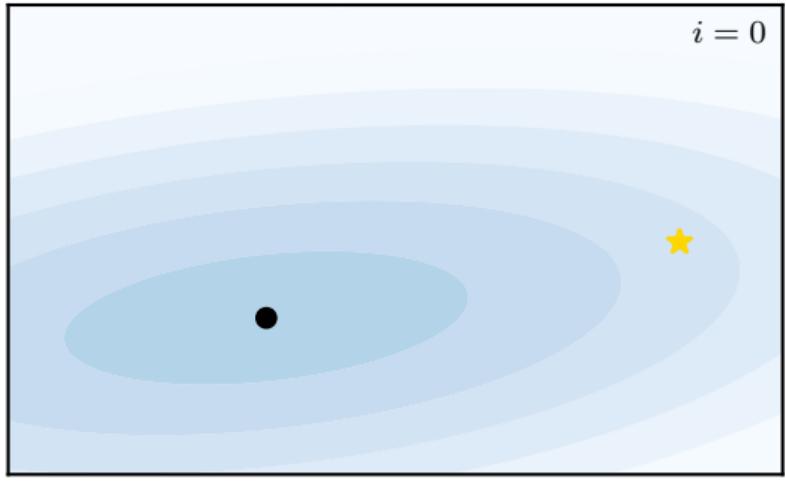
The Gaussian process prior makes assumptions about the representer weights.

[Wen+22]

## Observation:

GP prior induces representer weights prior:

$$\mathbf{y} - \boldsymbol{\mu} \sim \mathcal{N}\left(\mathbf{0}, \hat{\boldsymbol{\kappa}}\right)$$



- ★ Solution  $\mathbf{x}_*$
- Approximation  $\mathbf{x}_i$
- Belief  $p(\mathbf{x}_*)$

# Linear Solver Prior for GP Inference

The Gaussian process prior makes assumptions about the representer weights.

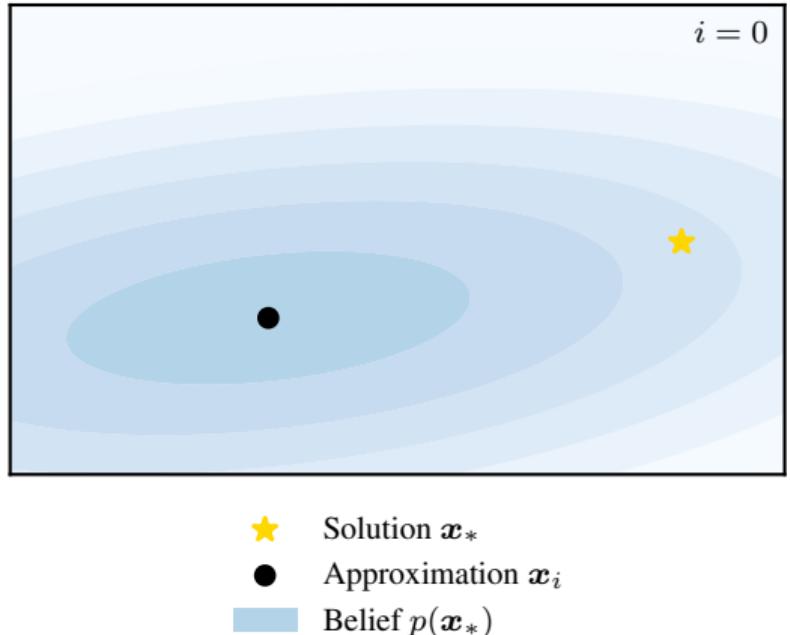
[Wen+22]

## Observation:

GP prior induces representer weights prior:

$$\mathbf{y} - \boldsymbol{\mu} \sim \mathcal{N}\left(\mathbf{0}, \hat{\boldsymbol{\kappa}}\right)$$

$$\Rightarrow \mathbf{v}_* = \hat{\boldsymbol{\kappa}}^{-1}(\mathbf{y} - \boldsymbol{\mu}) \sim \mathcal{N}\left(\begin{smallmatrix} \mathbf{0} \\ = \mathbf{v}_0 \end{smallmatrix}, \begin{smallmatrix} \hat{\boldsymbol{\kappa}}^{-1} \\ = \boldsymbol{\Sigma}_0 \end{smallmatrix}\right)$$





# Linear Solver Posterior for GP Inference

Estimation of representer weights with a probabilistic linear solver.

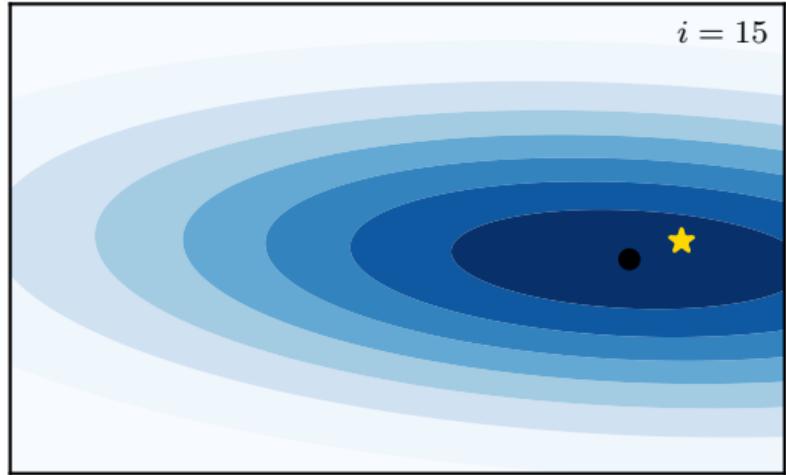
[Wen+22]

Representer weights posterior  $v_* \sim \mathcal{N}(v_i, \Sigma_i)$ , s.t.

$$v_i = \mathcal{C}_i(y - \mu)$$

$$\Sigma_i = \hat{K}^{-1} - \mathcal{C}_i$$

$$\hat{K}^{-1} = \begin{matrix} \text{[Pixelated matrix]} \end{matrix} \approx \begin{matrix} \text{[Pixelated matrix]} \end{matrix} = \mathcal{C}_i$$



- ★ Solution  $\mathbf{x}_*$
- Approximation  $\mathbf{x}_i$
- Belief  $p(\mathbf{x}_*)$



# Linear Solver Posterior for GP Inference

Estimation of representer weights with a probabilistic linear solver.

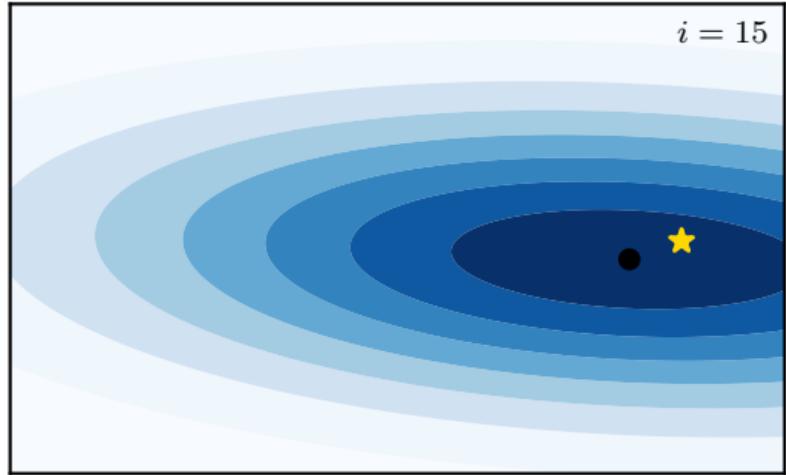
[Wen+22]

Representer weights posterior  $v_* \sim \mathcal{N}(v_i, \Sigma_i)$ , s.t.

$$v_i = \mathcal{C}_i(y - \mu)$$

$$\Sigma_i = \hat{K}^{-1} - \mathcal{C}_i$$

$$\hat{K}^{-1} = \begin{matrix} \text{[Pixelated matrix]} \end{matrix} \approx \begin{matrix} \text{[Pixelated matrix]} \end{matrix} = \mathcal{C}_i$$



- ★ Solution  $\mathbf{x}_*$
- Approximation  $\mathbf{x}_i$
- Belief  $p(\mathbf{x}_*)$

**Chicken & Egg Problem:** How can we get a probabilistic error estimate for  $v_i \approx v_*$ , if we need  $\hat{K}^{-1}$ ?

# IterGP: Computation-Aware Gaussian Process Inference



Quantifying uncertainty arising from observing finite data and performing a finite amount of computation.

[Wen+22]

**Goal:** Approximate the Gaussian process posterior  $f \mid \mathbf{y} \sim \mathcal{GP}(\mu_*, k_*)$ .

# IterGP: Computation-Aware Gaussian Process Inference



Quantifying uncertainty arising from observing finite data and performing a finite amount of computation.

[Wen+22]

**Goal:** Approximate the Gaussian process posterior  $f \mid \mathbf{y} \sim \mathcal{GP}(\mu_*, k_*)$ .

**Obtained:** Belief about representer weights  $\mathbf{v}_* = \hat{\mathbf{K}}^{-1}(\mathbf{y} - \boldsymbol{\mu}) \sim \mathcal{N}(\mathbf{v}_i, \mathbf{\Sigma}_{\mathbf{i}}) = \mathcal{N}\left(\mathbf{v}_i, \hat{\mathbf{K}}^{-1} - \mathbf{C}_i\right)$

# IterGP: Computation-Aware Gaussian Process Inference



Quantifying uncertainty arising from observing finite data and performing a finite amount of computation.

[Wen+22]

**Goal:** Approximate the Gaussian process posterior  $f \mid \mathbf{y} \sim \mathcal{GP}(\mu_*, k_*)$ .

**Obtained:** Belief about representer weights  $\mathbf{v}_* = \hat{\mathbf{K}}^{-1}(\mathbf{y} - \boldsymbol{\mu}) \sim \mathcal{N}(\mathbf{v}_i, \mathbf{\Sigma}_{\mathbf{i}}) = \mathcal{N}\left(\mathbf{v}_i, \hat{\mathbf{K}}^{-1} - \mathbf{C}_i\right)$

**Idea:** Propagate uncertainty about representer weights to posterior.



# IterGP: Computation-Aware Gaussian Process Inference

Quantifying uncertainty arising from observing finite data and performing a finite amount of computation.

[Wen+22]

**Goal:** Approximate the Gaussian process posterior  $f \mid \mathbf{y} \sim \mathcal{GP}(\mu_*, k_*)$ .

**Obtained:** Belief about representer weights  $\mathbf{v}_* = \hat{\mathbf{K}}^{-1}(\mathbf{y} - \boldsymbol{\mu}) \sim \mathcal{N}(\mathbf{v}_i, \mathbf{\Sigma}_i) = \mathcal{N}\left(\mathbf{v}_i, \hat{\mathbf{K}}^{-1} - \mathbf{C}_i\right)$

**Idea:** Propagate uncertainty about representer weights to posterior.

1 Pathwise form of posterior:  $(f \mid \mathbf{y})(\cdot) \stackrel{d}{=} f(\cdot) + k(\cdot, \mathbf{X}) \underbrace{\hat{\mathbf{K}}^{-1}(\mathbf{y} - \boldsymbol{\mu})}_{= \mathbf{v}_*} \stackrel{d}{=} (f \mid \mathbf{v}_*)(\cdot)$



# IterGP: Computation-Aware Gaussian Process Inference

Quantifying uncertainty arising from observing finite data and performing a finite amount of computation.

[Wen+22]

**Goal:** Approximate the Gaussian process posterior  $f \mid \mathbf{y} \sim \mathcal{GP}(\mu_*, k_*)$ .

**Obtained:** Belief about representer weights  $\mathbf{v}_* = \hat{\mathbf{K}}^{-1}(\mathbf{y} - \boldsymbol{\mu}) \sim \mathcal{N}(\mathbf{v}_i, \boldsymbol{\Sigma}_i) = \mathcal{N}\left(\mathbf{v}_i, \hat{\mathbf{K}}^{-1} - \mathbf{C}_i\right)$

**Idea:** Propagate uncertainty about representer weights to posterior.

1 Pathwise form of posterior:  $(f \mid \mathbf{y})(\cdot) \stackrel{d}{=} f(\cdot) + k(\cdot, \mathbf{X}) \underbrace{\hat{\mathbf{K}}^{-1}(\mathbf{y} - \boldsymbol{\mu})}_{= \mathbf{v}_*} \stackrel{d}{=} (f \mid \mathbf{v}_*)(\cdot)$



# IterGP: Computation-Aware Gaussian Process Inference

Quantifying uncertainty arising from observing finite data and performing a finite amount of computation.

[Wen+22]

**Goal:** Approximate the Gaussian process posterior  $f \mid \mathbf{y} \sim \mathcal{GP}(\mu_*, k_*)$ .

**Obtained:** Belief about representer weights  $\mathbf{v}_* = \hat{\mathbf{K}}^{-1}(\mathbf{y} - \boldsymbol{\mu}) \sim \mathcal{N}(\mathbf{v}_i, \boldsymbol{\Sigma}_i) = \mathcal{N}\left(\mathbf{v}_i, \hat{\mathbf{K}}^{-1} - \mathbf{C}_i\right)$

**Idea:** Propagate uncertainty about representer weights to posterior.

- 1 Pathwise form of posterior:  $(f \mid \mathbf{y})(\cdot) \stackrel{d}{=} f(\cdot) + k(\cdot, \mathbf{X}) \underbrace{\hat{\mathbf{K}}^{-1}(\mathbf{y} - \boldsymbol{\mu})}_{=\mathbf{v}_*} \stackrel{d}{=} (f \mid \mathbf{v}_*)(\cdot)$
- 2 Marginalize representer weights belief:  $p(f(\cdot)) = \int p(f(\cdot) \mid \mathbf{v}_*) p(\mathbf{v}_*) d\mathbf{v}_* = \mathcal{GP}(f; \mu_i, k_i)$



# IterGP: Computation-Aware Gaussian Process Inference

Quantifying uncertainty arising from observing finite data and performing a finite amount of computation.

[Wen+22]

**Goal:** Approximate the Gaussian process posterior  $f \mid \mathbf{y} \sim \mathcal{GP}(\mu_*, k_*)$ .

**Obtained:** Belief about representer weights  $\mathbf{v}_* = \hat{\mathbf{K}}^{-1}(\mathbf{y} - \boldsymbol{\mu}) \sim \mathcal{N}(\mathbf{v}_i, \boldsymbol{\Sigma}_i) = \mathcal{N}\left(\mathbf{v}_i, \hat{\mathbf{K}}^{-1} - \mathbf{C}_i\right)$

**Idea:** Propagate uncertainty about representer weights to posterior.

- 1 Pathwise form of posterior:  $(f \mid \mathbf{y})(\cdot) \stackrel{d}{=} f(\cdot) + k(\cdot, \mathbf{X}) \underbrace{\hat{\mathbf{K}}^{-1}(\mathbf{y} - \boldsymbol{\mu})}_{= \mathbf{v}_*} \stackrel{d}{=} (f \mid \mathbf{v}_*)(\cdot)$
- 2 Marginalize representer weights belief:  $p(f(\cdot)) = \int p(f(\cdot) \mid \mathbf{v}_*) p(\mathbf{v}_*) d\mathbf{v}_* = \mathcal{GP}(f; \mu_i, k_i)$

$$\mu_i(\cdot) = \mu(\cdot) + K(\cdot, \mathbf{X}) v_i$$



# IterGP: Computation-Aware Gaussian Process Inference

Quantifying uncertainty arising from observing finite data and performing a finite amount of computation.

[Wen+22]

**Goal:** Approximate the Gaussian process posterior  $f \mid \mathbf{y} \sim \mathcal{GP}(\mu_*, k_*)$ .

**Obtained:** Belief about representer weights  $\mathbf{v}_* = \hat{\mathbf{K}}^{-1}(\mathbf{y} - \boldsymbol{\mu}) \sim \mathcal{N}(\mathbf{v}_i, \mathbf{\Sigma}_i) = \mathcal{N}\left(\mathbf{v}_i, \hat{\mathbf{K}}^{-1} - \mathbf{C}_i\right)$

**Idea:** Propagate uncertainty about representer weights to posterior.

1 Pathwise form of posterior:  $(f \mid \mathbf{y})(\cdot) \stackrel{d}{=} f(\cdot) + k(\cdot, \mathbf{X}) \underbrace{\hat{\mathbf{K}}^{-1}(\mathbf{y} - \boldsymbol{\mu})}_{= \mathbf{v}_*} \stackrel{d}{=} (f \mid \mathbf{v}_*)(\cdot)$

2 Marginalize representer weights belief:  $p(f(\cdot)) = \int p(f(\cdot) \mid \mathbf{v}_*) p(\mathbf{v}_*) d\mathbf{v}_* = \mathcal{GP}(f; \mu_i, k_i)$

$$\mu_j(\cdot) = \mu(\cdot) + K(\cdot, \mathbf{X}) v_i$$

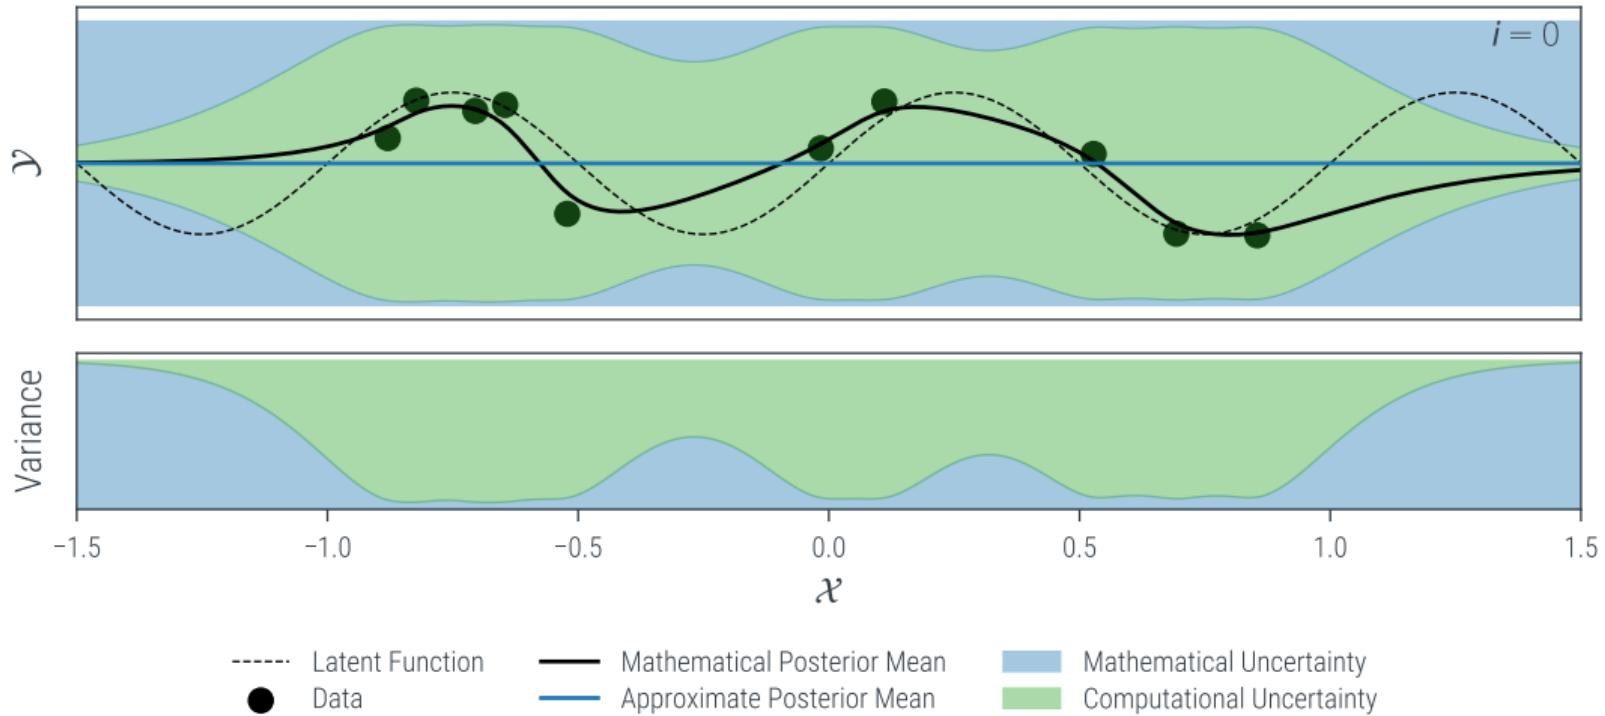
$$k_i(\cdot, \cdot) = \underbrace{K(\cdot, \cdot) - K(\cdot, \mathbf{X}) \hat{\mathbf{K}}^{-1} K(\mathbf{X}, \cdot)}_{= \mathbb{E}((f(\cdot) - \mu_*(\cdot))^2) \text{ mathematical uncertainty } \color{blue}{\bullet}} + \underbrace{K(\cdot, \mathbf{X}) \mathbf{\Sigma}_i K(\mathbf{X}, \cdot)}_{= \mathbb{E}((\mu_*(\cdot) - \mu_i(\cdot))^2) \text{ computational uncertainty } \color{green}{\bullet}} = \underbrace{K(\cdot, \cdot) - K(\cdot, \mathbf{X}) \mathbf{C}_i K(\mathbf{X}, \cdot)}_{= \mathbb{E}((f(\cdot) - \mu_i(\cdot))^2) \text{ combined uncertainty } \color{purple}{\bullet}}$$



# Computation-Aware GP Inference Illustrated

Interpreting computational and combined uncertainty as error quantification.

IterGP-PI

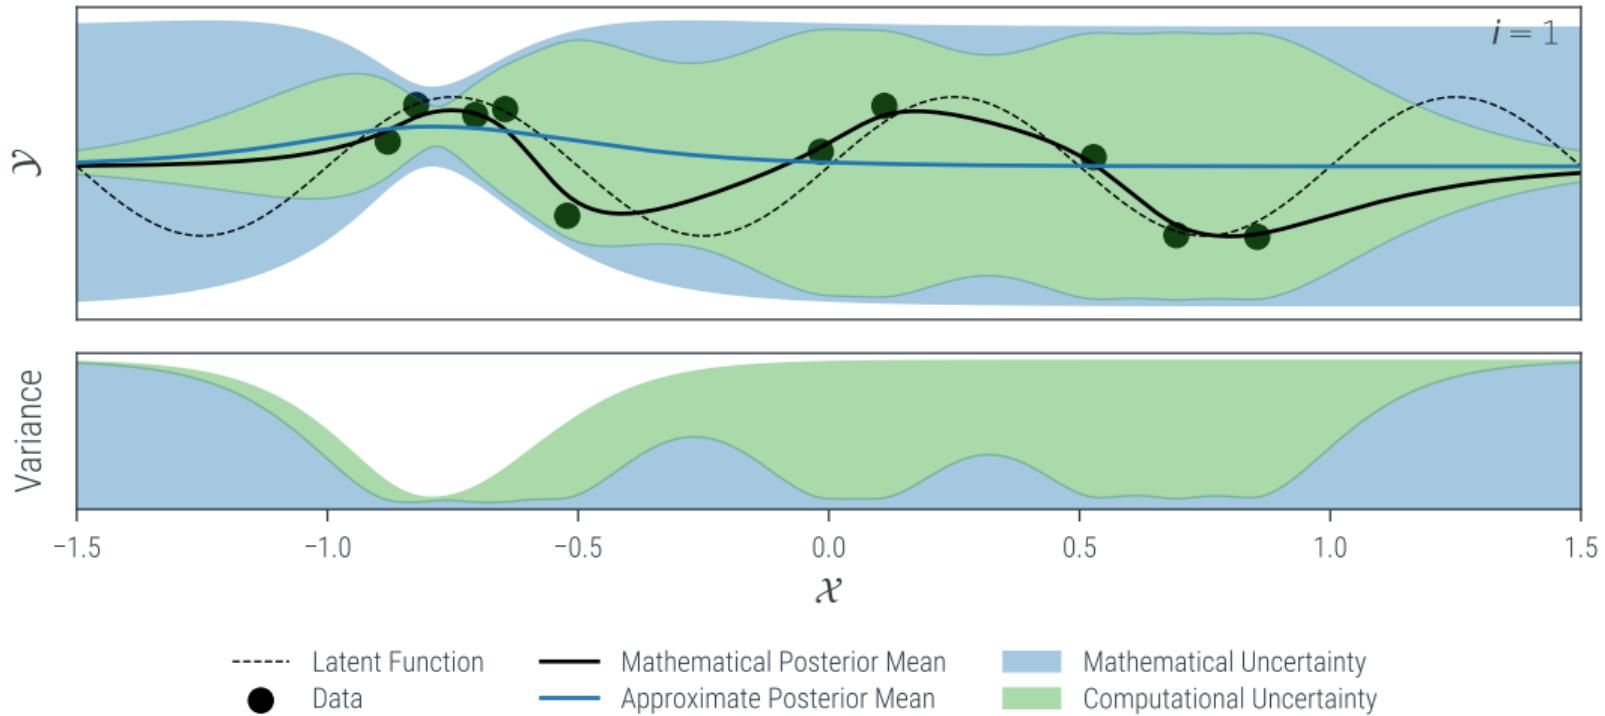




# Computation-Aware GP Inference Illustrated

Interpreting computational and combined uncertainty as error quantification.

IterGP-PI

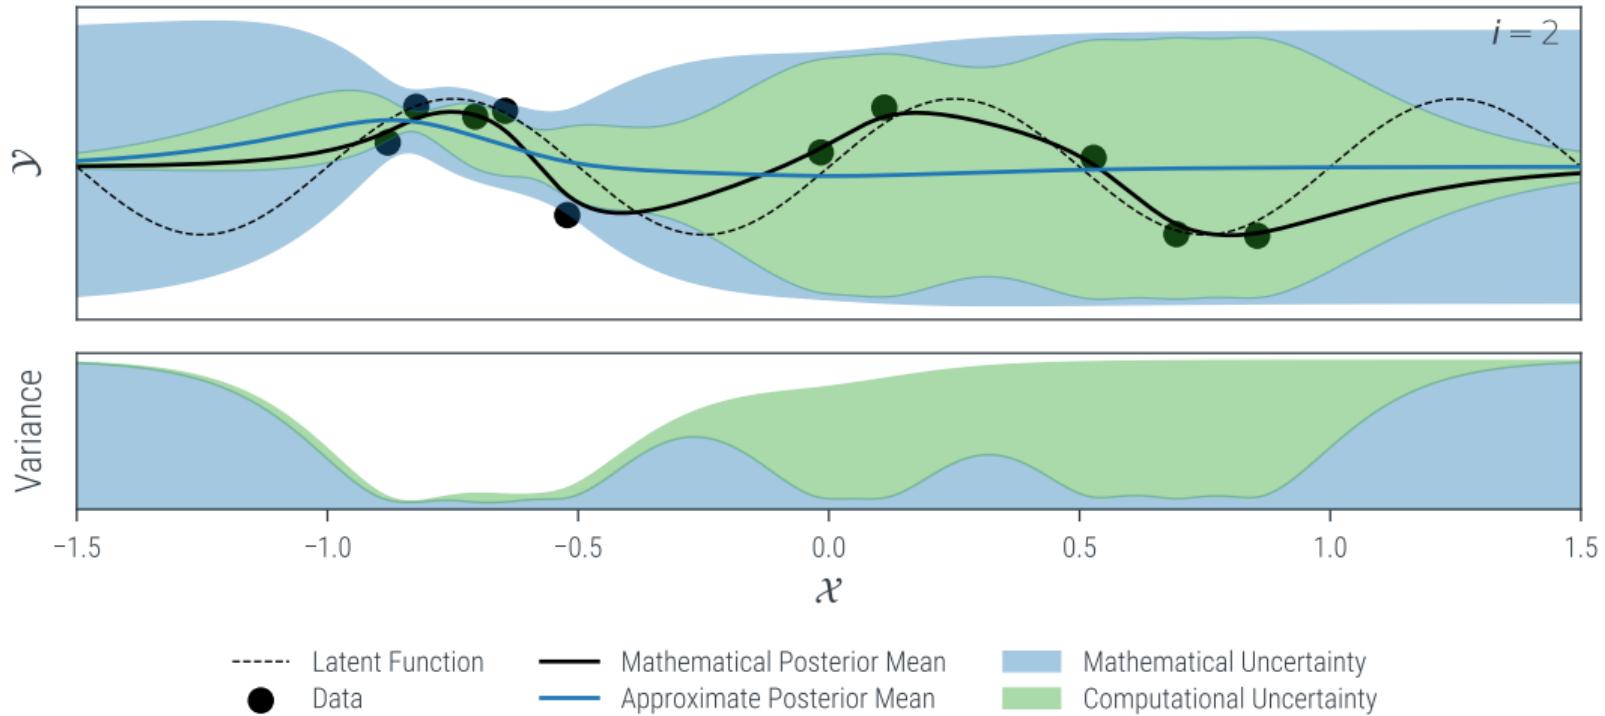




# Computation-Aware GP Inference Illustrated

Interpreting computational and combined uncertainty as error quantification.

IterGP-PI

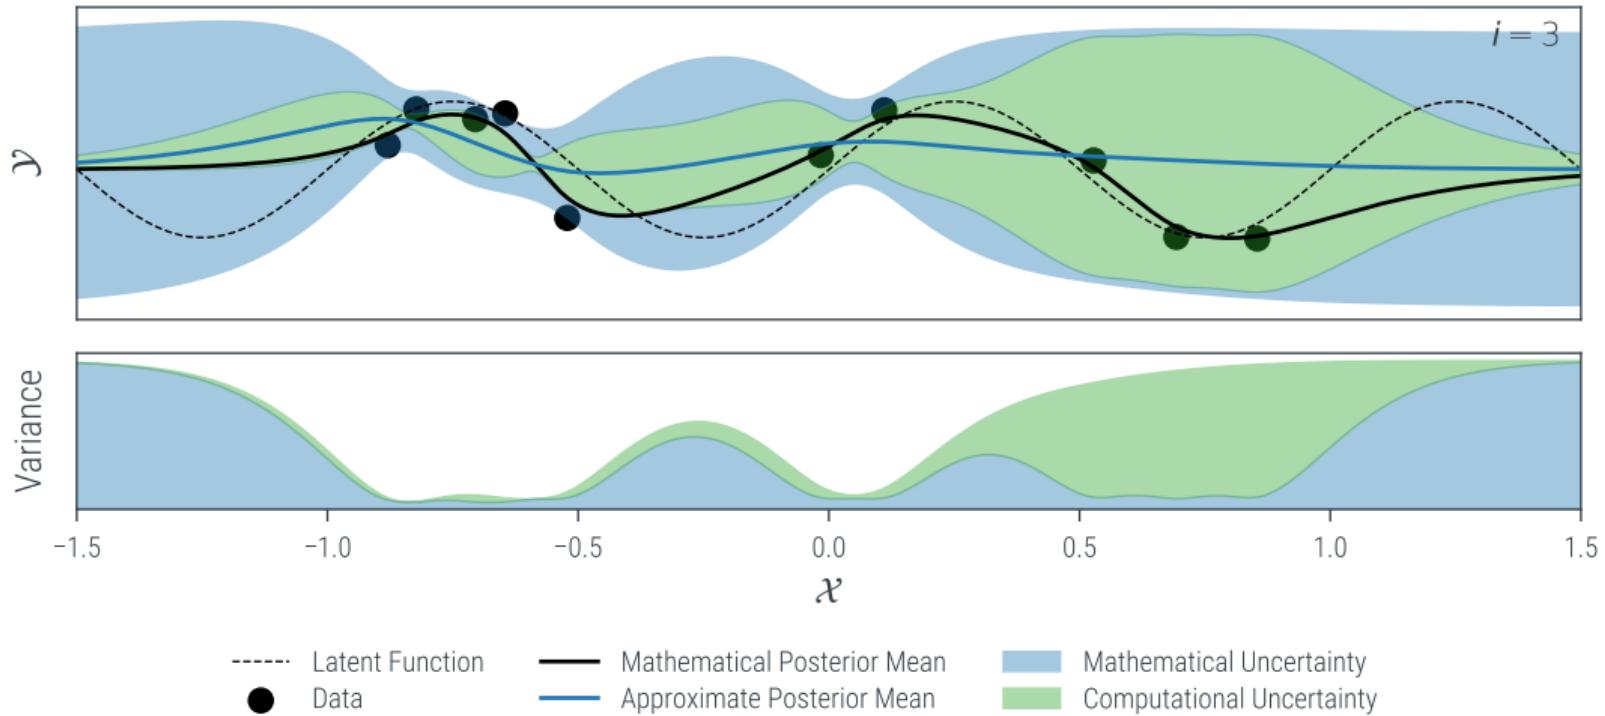




# Computation-Aware GP Inference Illustrated

Interpreting computational and combined uncertainty as error quantification.

IterGP-PI

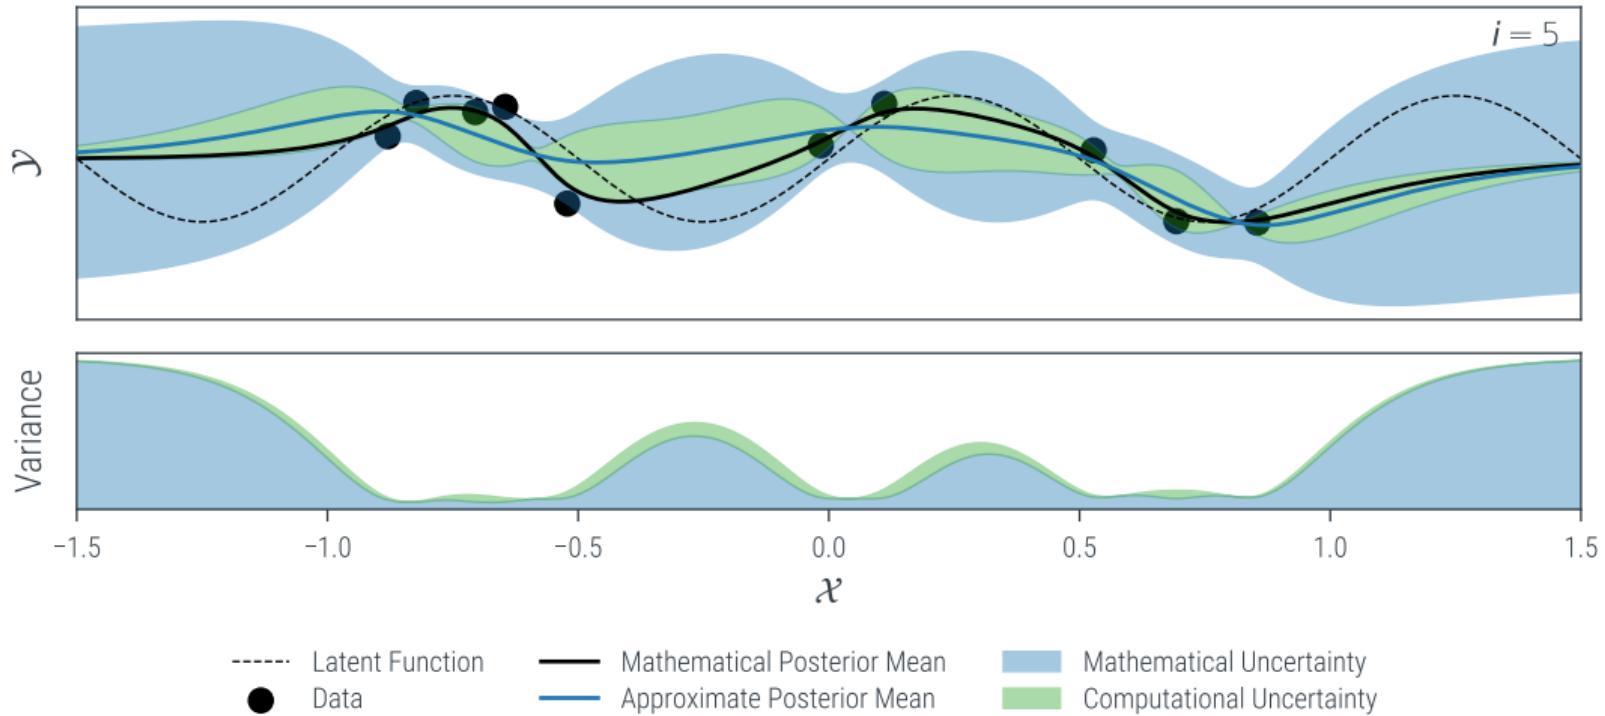




# Computation-Aware GP Inference Illustrated

Interpreting computational and combined uncertainty as error quantification.

IterGP-PI





# Algorithm: IterGP

◀ ToC

---

		Time	Space
1	procedure ITERGP( $\mu, K, X, y, C_0 = \mathbf{0}$ )		
2	while not STOPPINGCRITERION() do		
3	$s_i \leftarrow \text{POLICY}()$	Select action via policy.	
4	$r_{i-1} \leftarrow (y - \mu) - \hat{K}v_{i-1}$	Residual.	$\mathcal{O}(n^2)$
5	$\alpha_i \leftarrow s_i^\top r_{i-1}$	Observation.	$\mathcal{O}(n)$
6	$z_i \leftarrow \hat{K}s_i$		$\mathcal{O}(n^2)$
7	$d_i \leftarrow \Sigma_{i-1} \hat{K} s_i = s_i - C_{i-1} z_i$	Search direction.	$\mathcal{O}(ni)$
8	$\eta_i \leftarrow s_i^\top \hat{K} \Sigma_{i-1} \hat{K} s_i = z_i^\top d_i$		$\mathcal{O}(n)$
9	$C_i \leftarrow C_{i-1} + \frac{1}{\eta_i} d_i d_i^\top$	Precision matrix approx. $C_i \approx \hat{K}^{-1}$ .	$\mathcal{O}(n)$
10	$v_i \leftarrow v_{i-1} + \frac{\alpha_i}{\eta_i} d_i$	Representer weights estimate.	$\mathcal{O}(n)$
11	$\Sigma_i \leftarrow \Sigma_0 - C_i$	Representer weights uncertainty.	
12	$\mu_i(\cdot) \leftarrow \mu(\cdot) + K(\cdot, X)v_i$	Approximate posterior mean.	$\mathcal{O}(n_o n)$
13	$K_i(\cdot, \cdot) \leftarrow K(\cdot, \cdot) - K(\cdot, X)C_i K(X, \cdot)$	Combined covariance function.	$\mathcal{O}(n_o ni)$
14	return $\mathcal{GP}(\mu_i, K_i)$		$\mathcal{O}(n_o^2)$

---



# Theoretical Analysis

Uncertainty as a tight bound on the relative error.

Theorem (Kanagawa et al. [Kan+18])

$$\sup_{g \in \mathcal{H}_{k\sigma} : \|g\|_{\mathcal{H}_{k\sigma}} \leq 1} \underbrace{g(x) - \mu_*^g(x)}_{\text{error of math. post. mean}} = \sup_{g \in \mathcal{H}_{k\sigma}} \frac{|g(x) - \mu_*^g(x)|}{\|g\|_{\mathcal{H}_{k\sigma}}} = \sqrt{k_*(x, x) + \sigma^2}$$



# Theoretical Analysis

Uncertainty as a tight bound on the relative error.

Theorem (Kanagawa et al. [Kan+18])

$$\sup_{g \in \mathcal{H}_{k\sigma} : \|g\|_{\mathcal{H}_{k\sigma}} \leq 1} \underbrace{\frac{g(\mathbf{x}) - \mu_*^g(\mathbf{x})}{\text{error of math. post. mean } \bullet}}_{\text{error of approximate posterior mean } \bullet + \text{computational error } \bullet} = \sup_{g \in \mathcal{H}_{k\sigma}} \frac{|g(\mathbf{x}) - \mu_*^g(\mathbf{x})|}{\|g\|_{\mathcal{H}_{k\sigma}}} = \sqrt{k_*(\mathbf{x}, \mathbf{x}) + \sigma^2}$$

Theorem (Wenger et al. [Wen+22])

$$\sup_{g \in \mathcal{H}_{k\sigma} : \|g\|_{\mathcal{H}_{k\sigma}} \leq 1} \underbrace{\frac{g(\mathbf{x}) - \mu_*^g(\mathbf{x})}{\text{error of math. post. mean } \bullet}}_{\text{error of approximate posterior mean } \bullet + \text{computational error } \bullet} + \underbrace{\frac{\mu_*^g(\mathbf{x}) - \mu_i^g(\mathbf{x})}{\text{computational error } \bullet}}_{\text{computational error } \bullet} = \sqrt{k_i(\mathbf{x}, \mathbf{x}) + \sigma^2}$$



# Theoretical Analysis

Uncertainty as a tight bound on the relative error.

Theorem (Kanagawa et al. [Kan+18])

$$\sup_{g \in \mathcal{H}_{k\sigma} : \|g\|_{\mathcal{H}_{k\sigma}} \leq 1} \underbrace{\frac{g(\mathbf{x}) - \mu_*^g(\mathbf{x})}{\text{error of math. post. mean } \bullet}}_{\text{error of approximate posterior mean } \circlearrowleft} = \sup_{g \in \mathcal{H}_{k\sigma}} \frac{|g(\mathbf{x}) - \mu_*^g(\mathbf{x})|}{\|g\|_{\mathcal{H}_{k\sigma}}} = \sqrt{k_*(\mathbf{x}, \mathbf{x}) + \sigma^2}$$

Theorem (Wenger et al. [Wen+22])

$$\sup_{g \in \mathcal{H}_{k\sigma} : \|g\|_{\mathcal{H}_{k\sigma}} \leq 1} \underbrace{\frac{g(\mathbf{x}) - \mu_*^g(\mathbf{x})}{\text{error of math. post. mean } \bullet}}_{\text{error of approximate posterior mean } \circlearrowleft} + \underbrace{\frac{\mu_*^g(\mathbf{x}) - \mu_i^g(\mathbf{x})}{\text{computational error } \circlearrowright}}_{\text{computational error } \circlearrowright} = \sqrt{k_i(\mathbf{x}, \mathbf{x}) + \sigma^2}$$

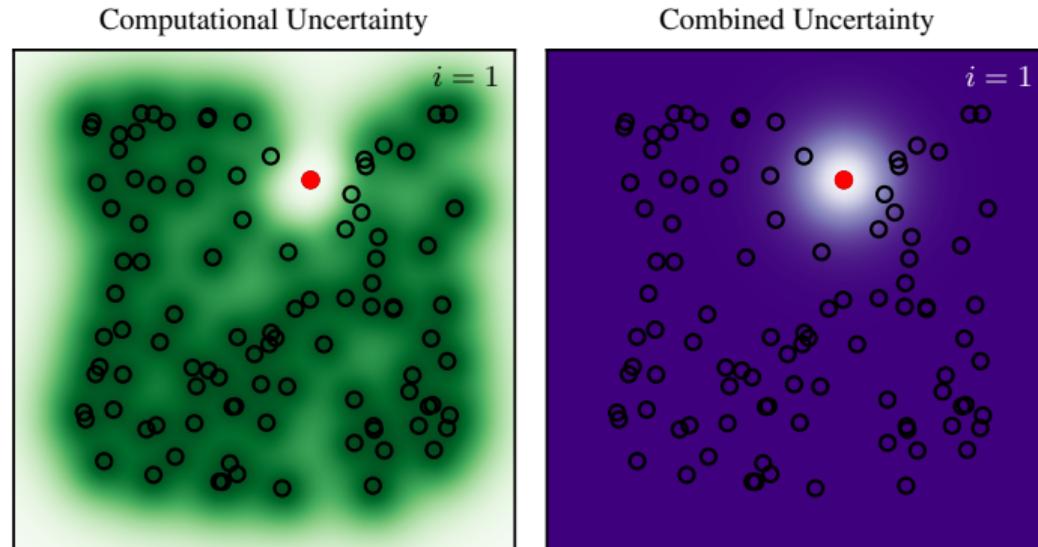
Exact quantification of uncertainty from **limited data** and **limited computation**.



# Policy Choice and Connection to Other Approximations

IterGP extends the most commonly used GP approximations to include computational uncertainty, with at most quadratic cost.

	Actions $s_i$	Classic Analog
IterGP-Cholesky	$e_i$	(Partial) Cholesky / Subset of data
IterGP-CG	$\hat{P}^{-1}r_i$	(Preconditioned) CG
IterGP-PseudoInput	$k(X, z_i)$	$\approx$ SVGP

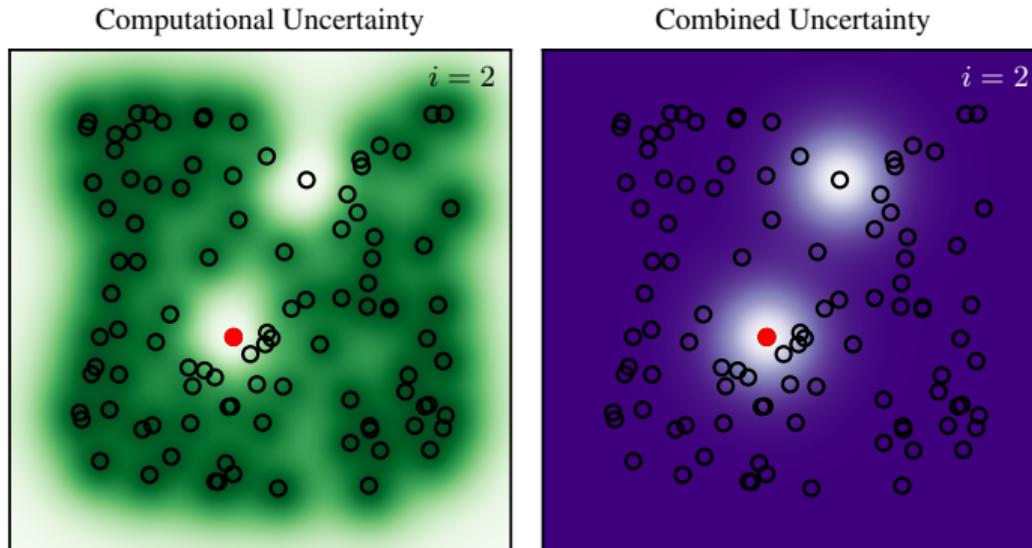




# Policy Choice and Connection to Other Approximations

IterGP extends the most commonly used GP approximations to include computational uncertainty, with at most quadratic cost.

	Actions $s_i$	Classic Analog
IterGP-Cholesky	$e_i$	(Partial) Cholesky / Subset of data
IterGP-CG	$\hat{P}^{-1}r_i$	(Preconditioned) CG
IterGP-PseudoInput	$k(X, z_i)$	$\approx$ SVGP

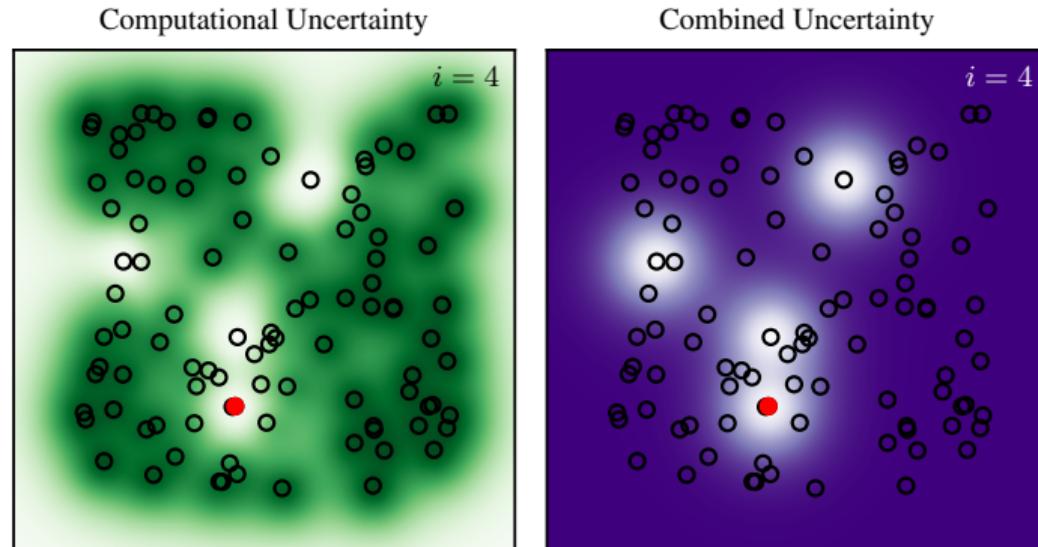




# Policy Choice and Connection to Other Approximations

IterGP extends the most commonly used GP approximations to include computational uncertainty, with at most quadratic cost.

	Actions $s_i$	Classic Analog
IterGP-Cholesky	$e_i$	(Partial) Cholesky / Subset of data
IterGP-CG	$\hat{P}^{-1}r_i$	(Preconditioned) CG
IterGP-PseudoInput	$k(X, z_i)$	$\approx$ SVGP

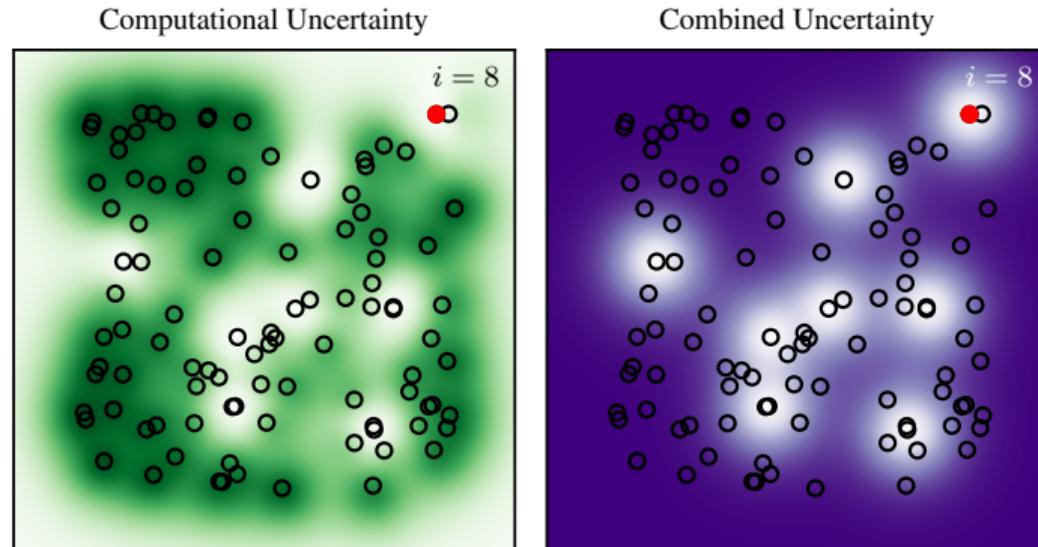




# Policy Choice and Connection to Other Approximations

IterGP extends the most commonly used GP approximations to include computational uncertainty, with at most quadratic cost.

	Actions $s_i$	Classic Analog
IterGP-Cholesky	$e_i$	(Partial) Cholesky / Subset of data
IterGP-CG	$\hat{P}^{-1}r_i$	(Preconditioned) CG
IterGP-PseudoInput	$k(X, z_i)$	$\approx$ SVGP

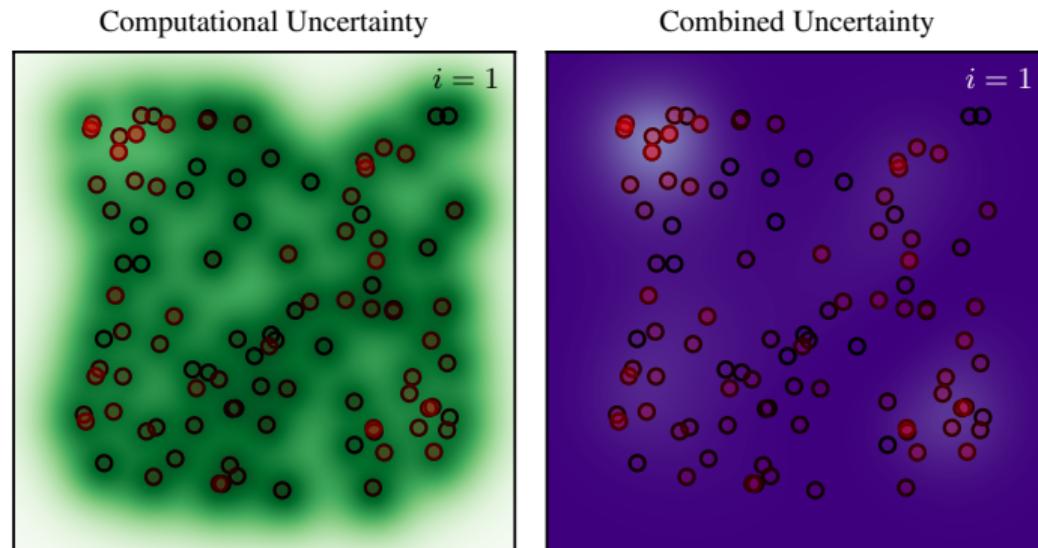




# Policy Choice and Connection to Other Approximations

IterGP extends the most commonly used GP approximations to include computational uncertainty, with at most quadratic cost.

	Actions $s_i$	Classic Analog
IterGP-Cholesky	$e_i$	(Partial) Cholesky / Subset of data
IterGP-CG	$\hat{P}^{-1}r_i$	(Preconditioned) CG
IterGP-PseudoInput	$k(X, z_i)$	$\approx$ SVGP

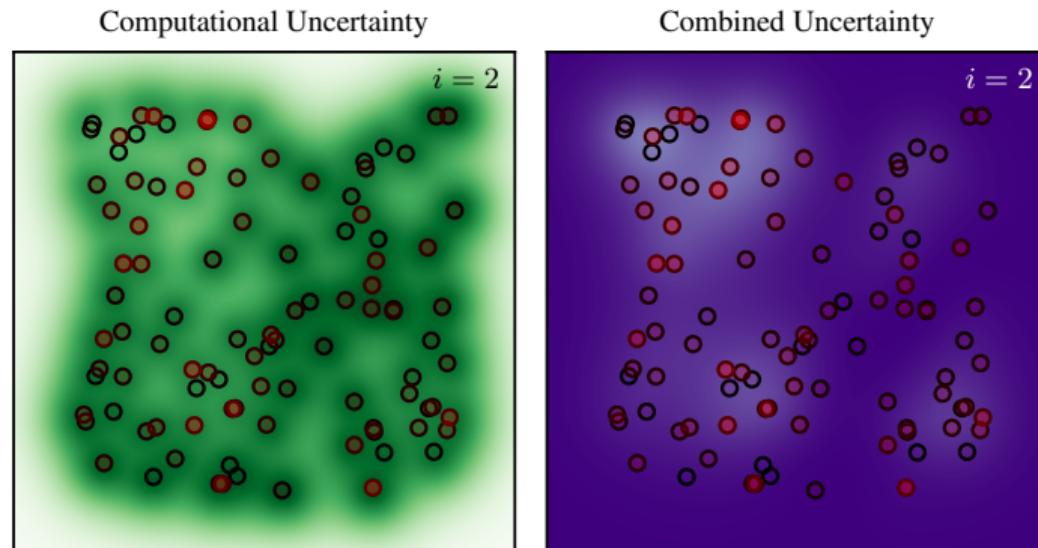




# Policy Choice and Connection to Other Approximations

IterGP extends the most commonly used GP approximations to include computational uncertainty, with at most quadratic cost.

	Actions $s_i$	Classic Analog
IterGP-Cholesky	$e_i$	(Partial) Cholesky / Subset of data
IterGP-CG	$\hat{P}^{-1}r_i$	(Preconditioned) CG
IterGP-PseudoInput	$k(X, z_i)$	$\approx$ SVGP

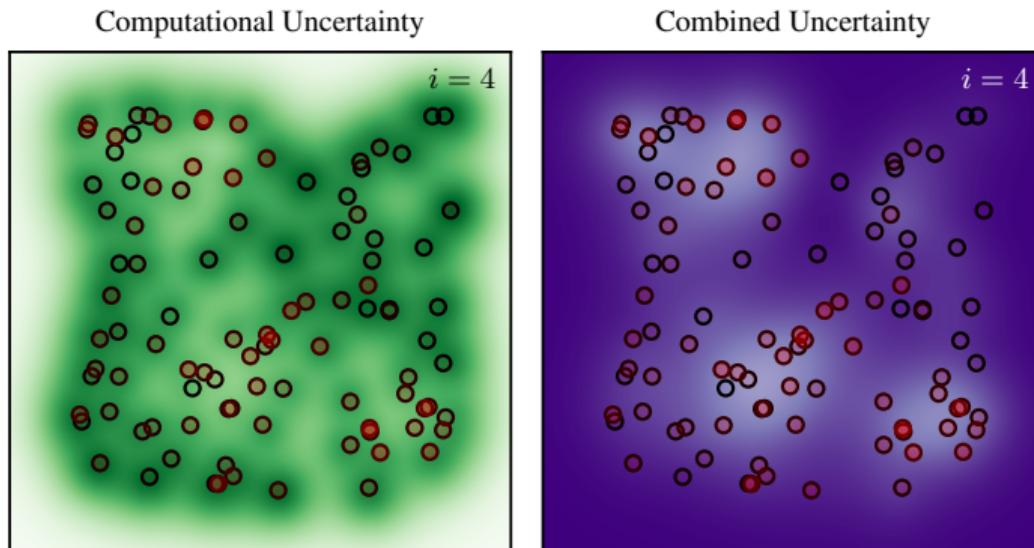




# Policy Choice and Connection to Other Approximations

IterGP extends the most commonly used GP approximations to include computational uncertainty, with at most quadratic cost.

	Actions $s_i$	Classic Analog
IterGP-Cholesky	$e_i$	(Partial) Cholesky / Subset of data
IterGP-CG	$\hat{P}^{-1}r_i$	(Preconditioned) CG
IterGP-PseudoInput	$k(X, z_i)$	$\approx$ SVGP

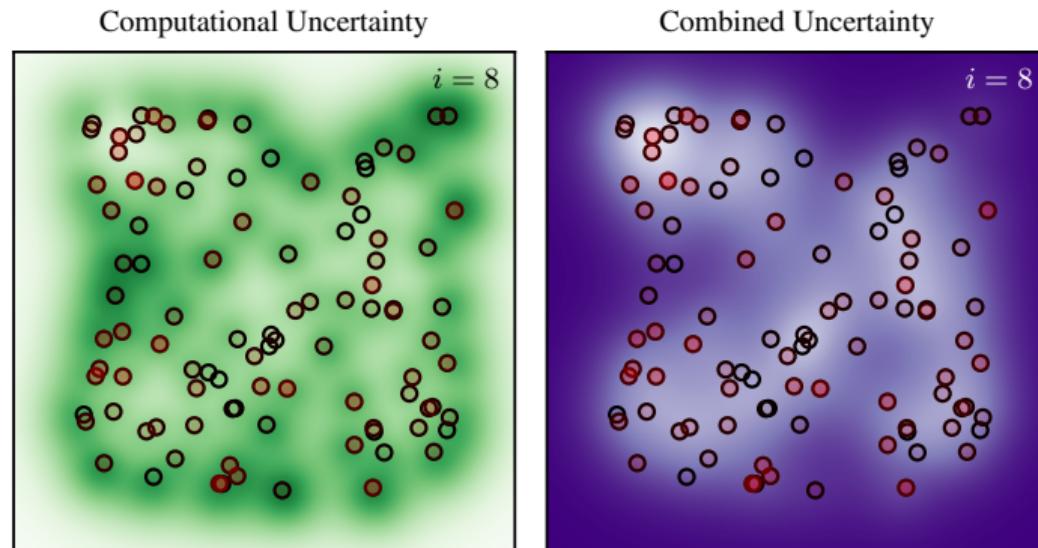




# Policy Choice and Connection to Other Approximations

IterGP extends the most commonly used GP approximations to include computational uncertainty, with at most quadratic cost.

	Actions $s_i$	Classic Analog
IterGP-Cholesky	$e_i$	(Partial) Cholesky / Subset of data
IterGP-CG	$\hat{P}^{-1}r_i$	(Preconditioned) CG
IterGP-Pseudoinput	$k(X, z_i)$	$\approx$ SVGP

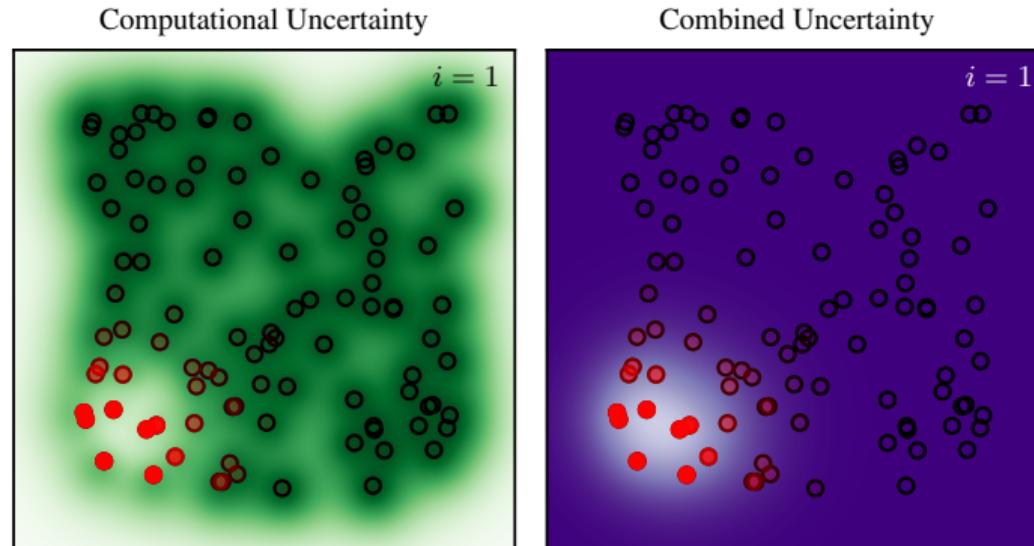




# Policy Choice and Connection to Other Approximations

IterGP extends the most commonly used GP approximations to include computational uncertainty, with at most quadratic cost.

	Actions $s_i$	Classic Analog
IterGP-Cholesky	$e_i$	(Partial) Cholesky / Subset of data
IterGP-CG	$\hat{P}^{-1}r_i$	(Preconditioned) CG
IterGP-Pseudoinput	$k(X, z_i)$	$\approx$ SVGP

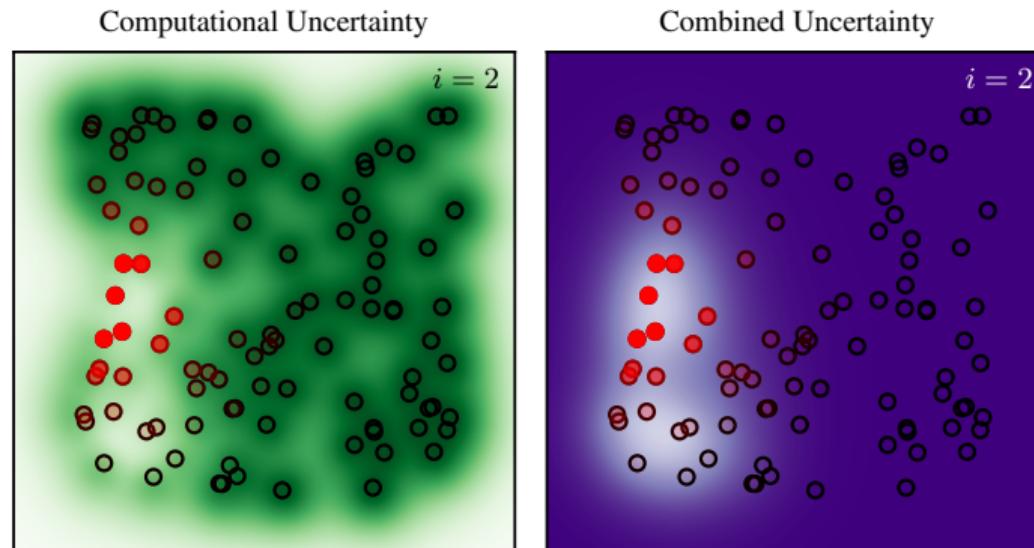




# Policy Choice and Connection to Other Approximations

IterGP extends the most commonly used GP approximations to include computational uncertainty, with at most quadratic cost.

	Actions $s_i$	Classic Analog
IterGP-Cholesky	$e_i$	(Partial) Cholesky / Subset of data
IterGP-CG	$\hat{P}^{-1}r_i$	(Preconditioned) CG
IterGP-Pseudoinput	$k(X, z_i)$	$\approx$ SVGP

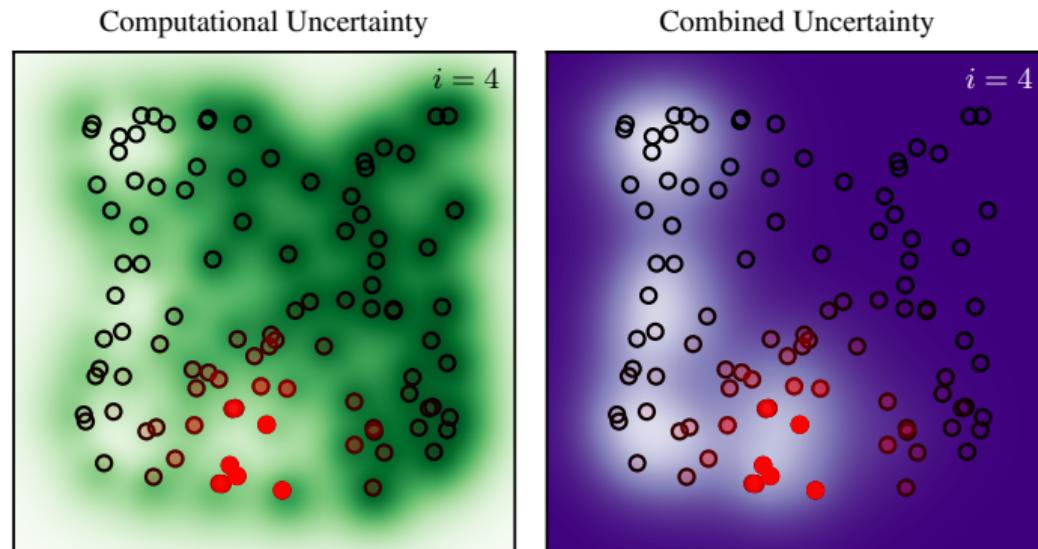




# Policy Choice and Connection to Other Approximations

IterGP extends the most commonly used GP approximations to include computational uncertainty, with at most quadratic cost.

	Actions $s_i$	Classic Analog
IterGP-Cholesky	$e_i$	(Partial) Cholesky / Subset of data
IterGP-CG	$\hat{P}^{-1}r_i$	(Preconditioned) CG
IterGP-Pseudoinput	$k(X, z_i)$	$\approx$ SVGP

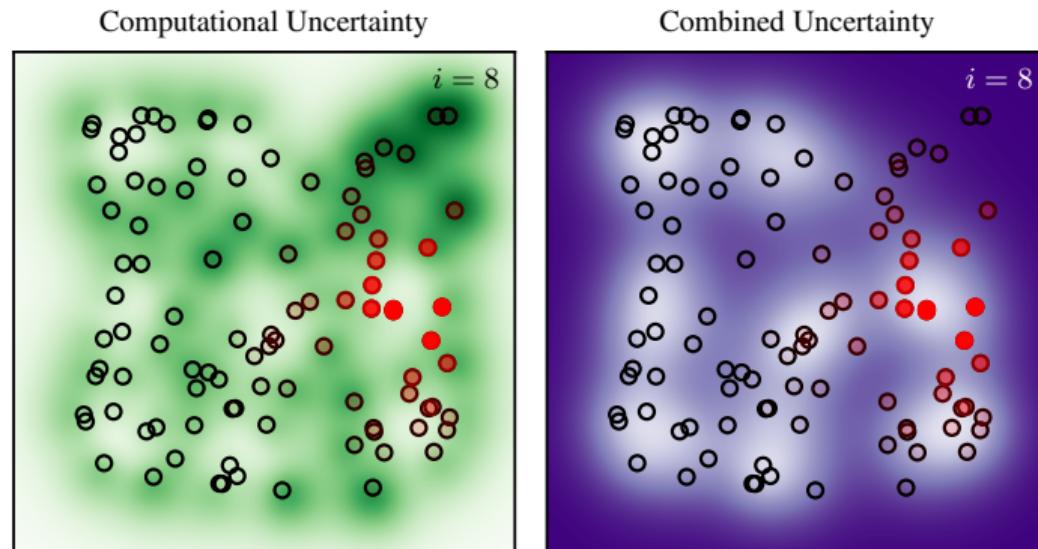




# Policy Choice and Connection to Other Approximations

IterGP extends the most commonly used GP approximations to include computational uncertainty, with at most quadratic cost.

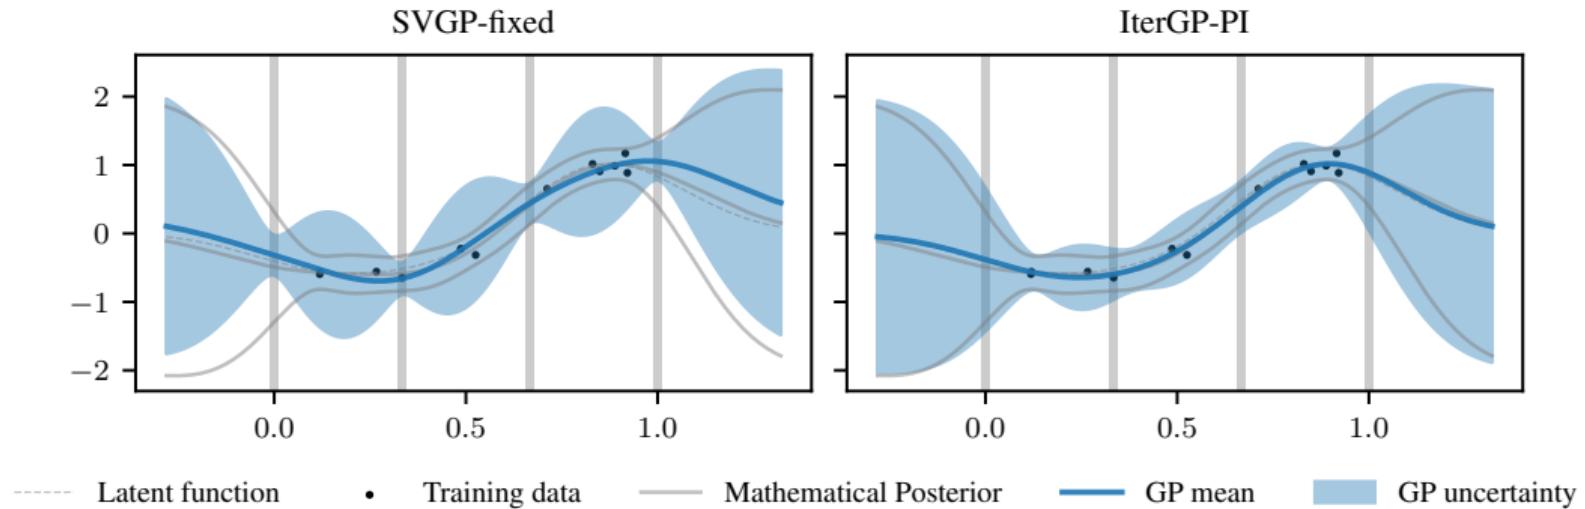
	Actions $s_i$	Classic Analog
IterGP-Cholesky	$e_i$	(Partial) Cholesky / Subset of data
IterGP-CG	$\hat{P}^{-1}r_i$	(Preconditioned) CG
IterGP-Pseudoinput	$k(X, z_i)$	$\approx$ SVGP



# SVGP versus IterGP-PI

Quantifying computational uncertainty improves generalization of inducing point methods like SVGP.

[Tit09; HFL13]

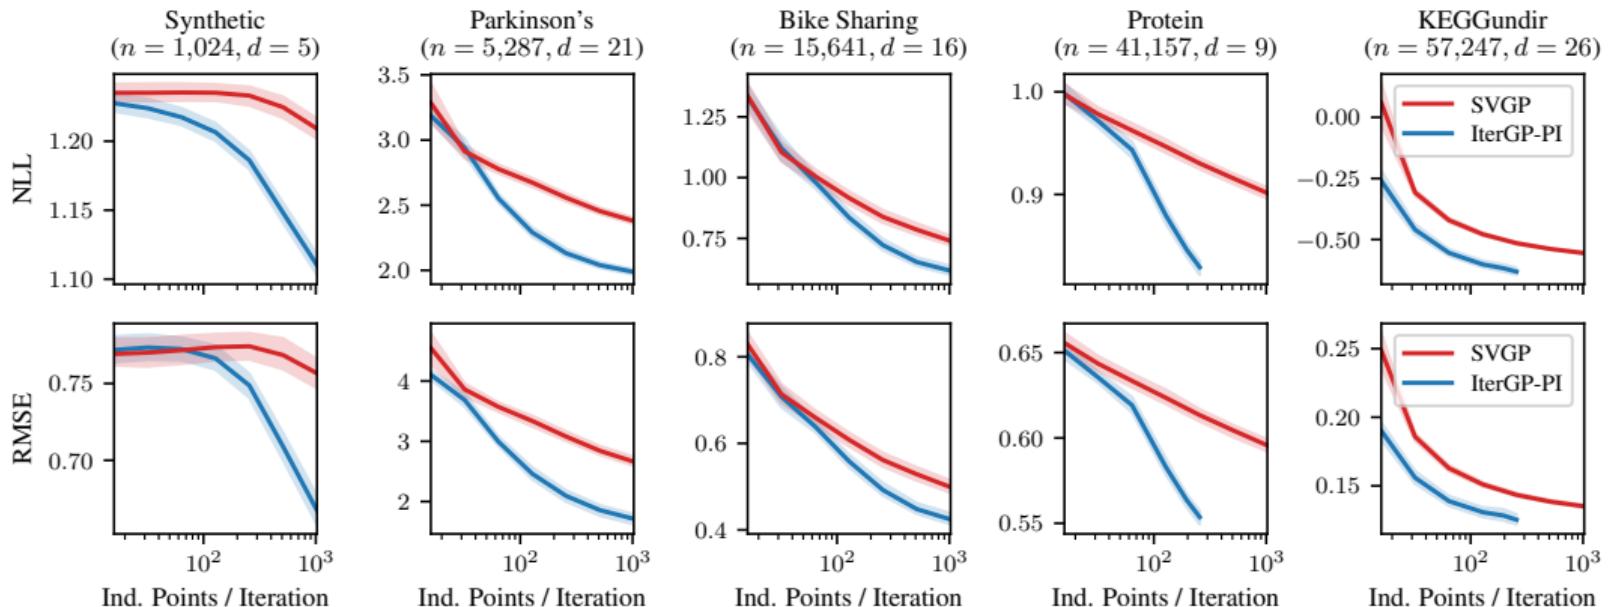




# SVGP versus IterGP-PI

Quantifying computational uncertainty improves generalization of inducing point methods like SVGP.

[Tit09; HFL13]

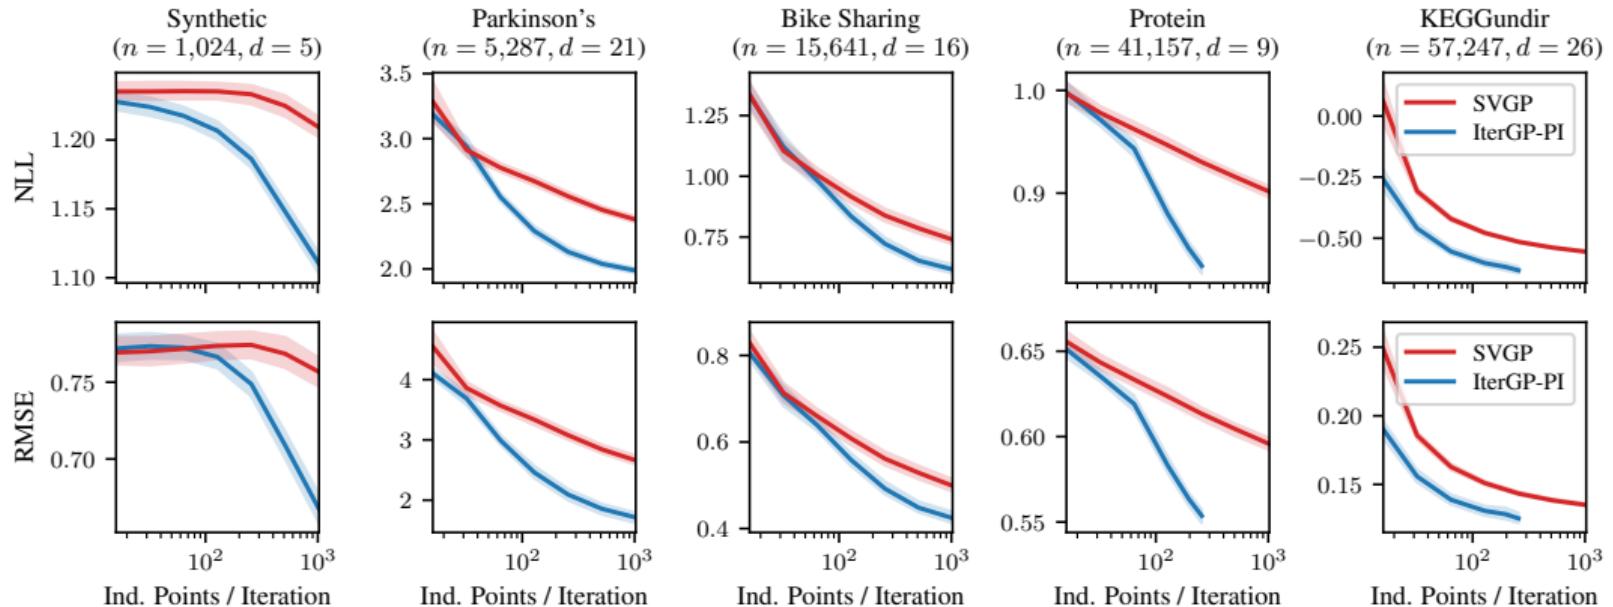




# SVGP versus IterGP-PI

Quantifying computational uncertainty improves generalization of inducing point methods like SVGP.

[Tit09; HFL13]

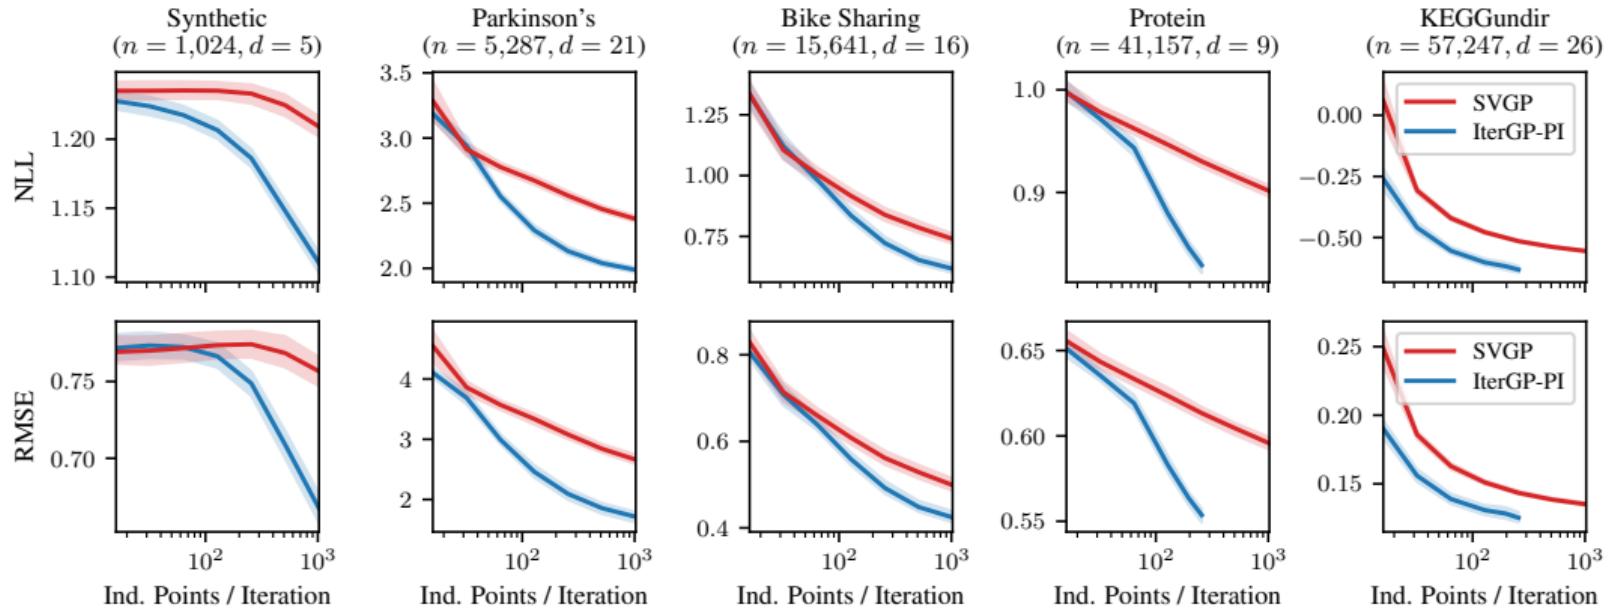


What about optimizing inducing point locations?

# SVGP versus IterGP-PI

Quantifying computational uncertainty improves generalization of inducing point methods like SVGP.

[Tit09; HFL13]



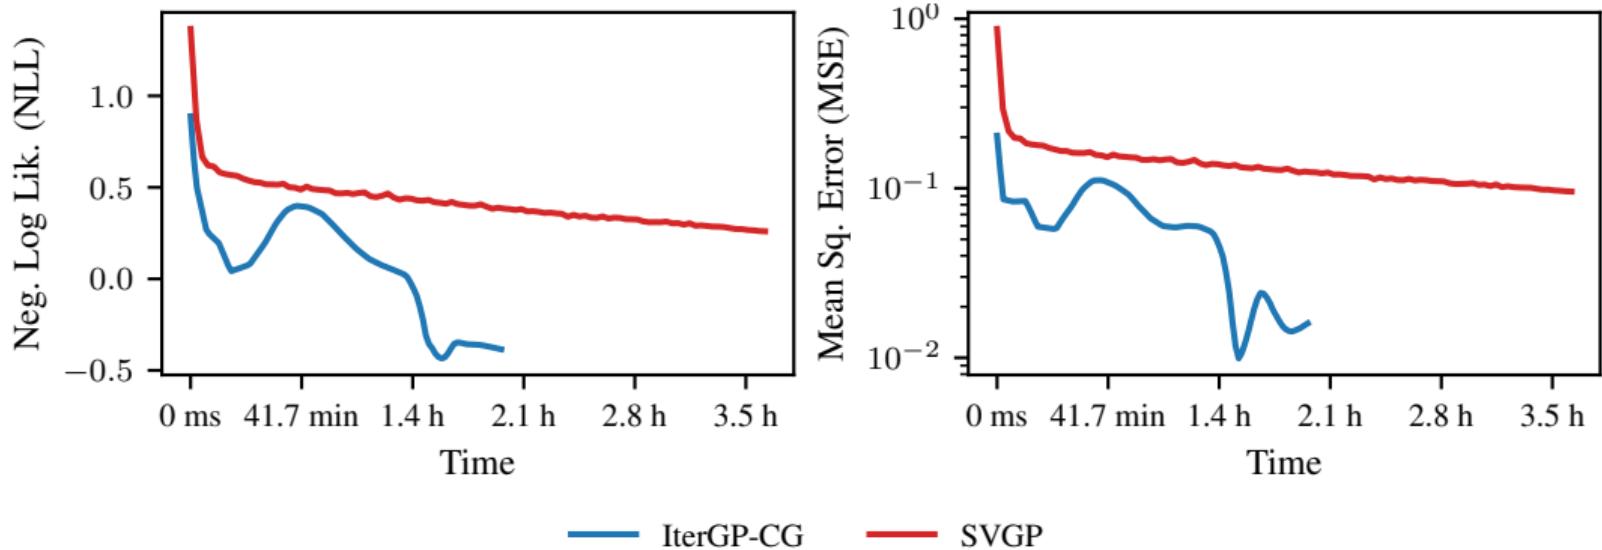
What about computational cost? SVGP:  $\mathcal{O}(ni^2)$  versus IterGP:  $\mathcal{O}(n^2i)$ .



# Training Gaussian Processes on Large-Scale Data

Kernel hyperparameter optimization with SVGP and IterGP on a problem with  $n \approx 4 \cdot 10^5$  data points.

[Wen+24, Unpublished work]

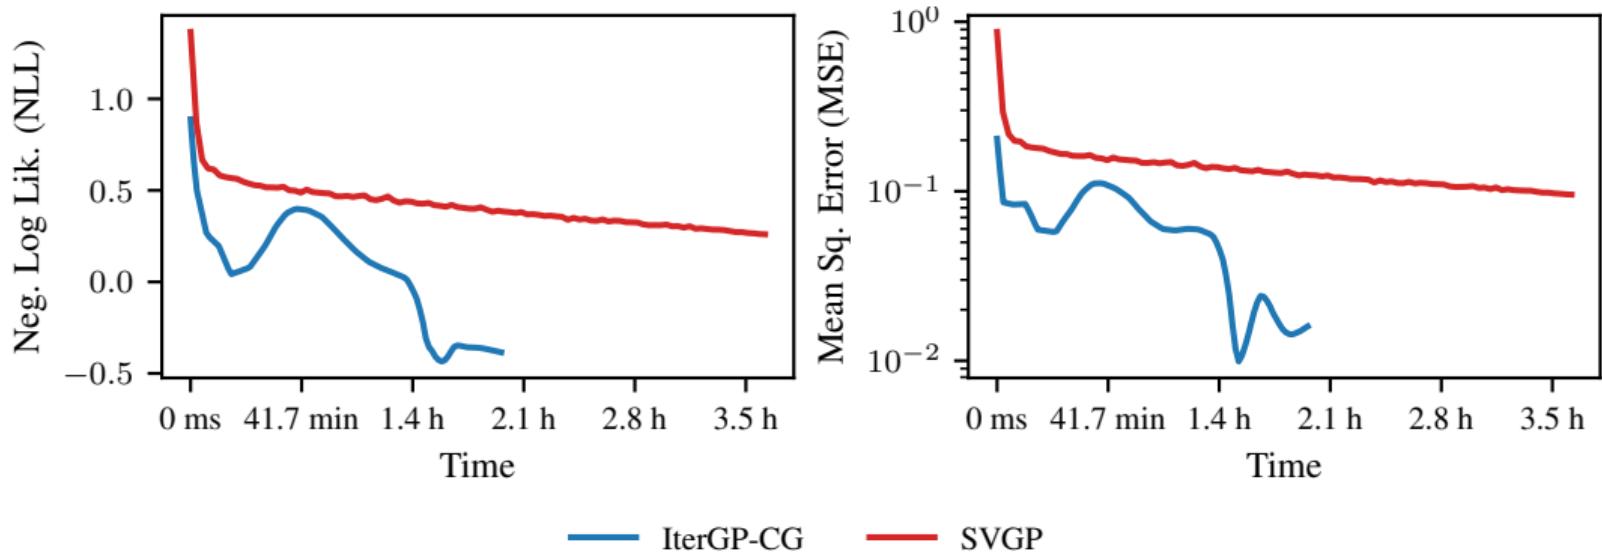




# Training Gaussian Processes on Large-Scale Data

Kernel hyperparameter optimization with SVGP and IterGP on a problem with  $n \approx 4 \cdot 10^5$  data points.

[Wen+24, Unpublished work]



Faster large-scale Gaussian processes with better generalization!

# Other Applications

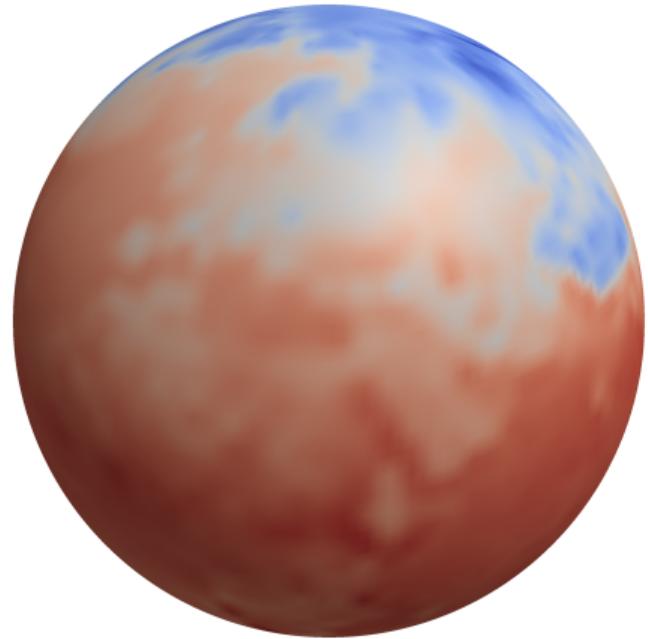
Extending these ideas beyond what we've seen.



# Spatiotemporal Modeling

Spatio-temporal regression of Earth surface temperature via computation-aware filtering and smoothing.

[Pfö+24, Unpublished Work]



(a) Prediction



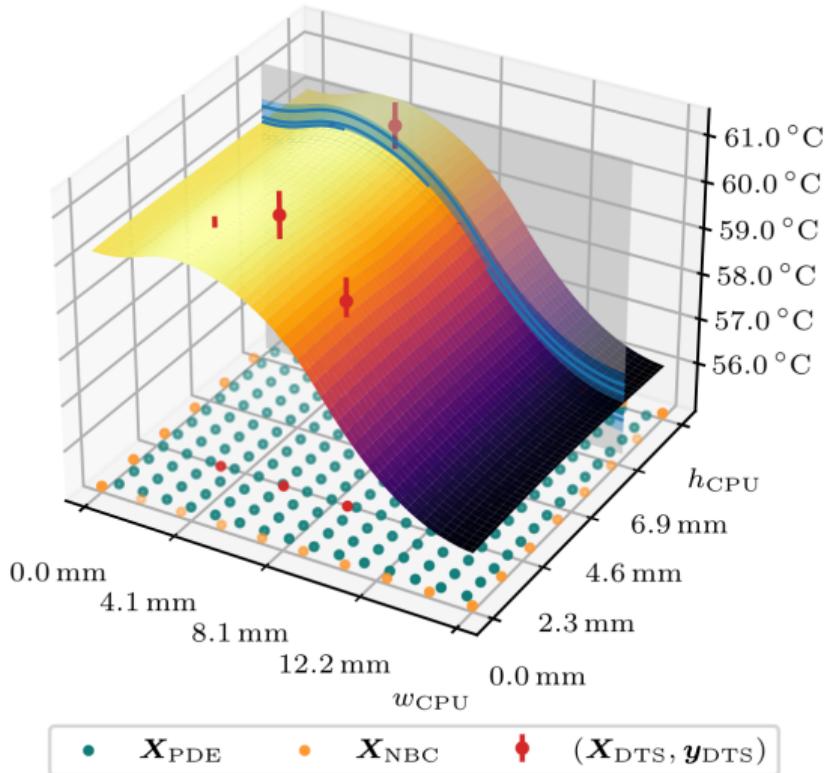
(b) Uncertainty



# Physics-Informed GP Regression

Learning to solve linear partial differential equations.

[Pfö+23]

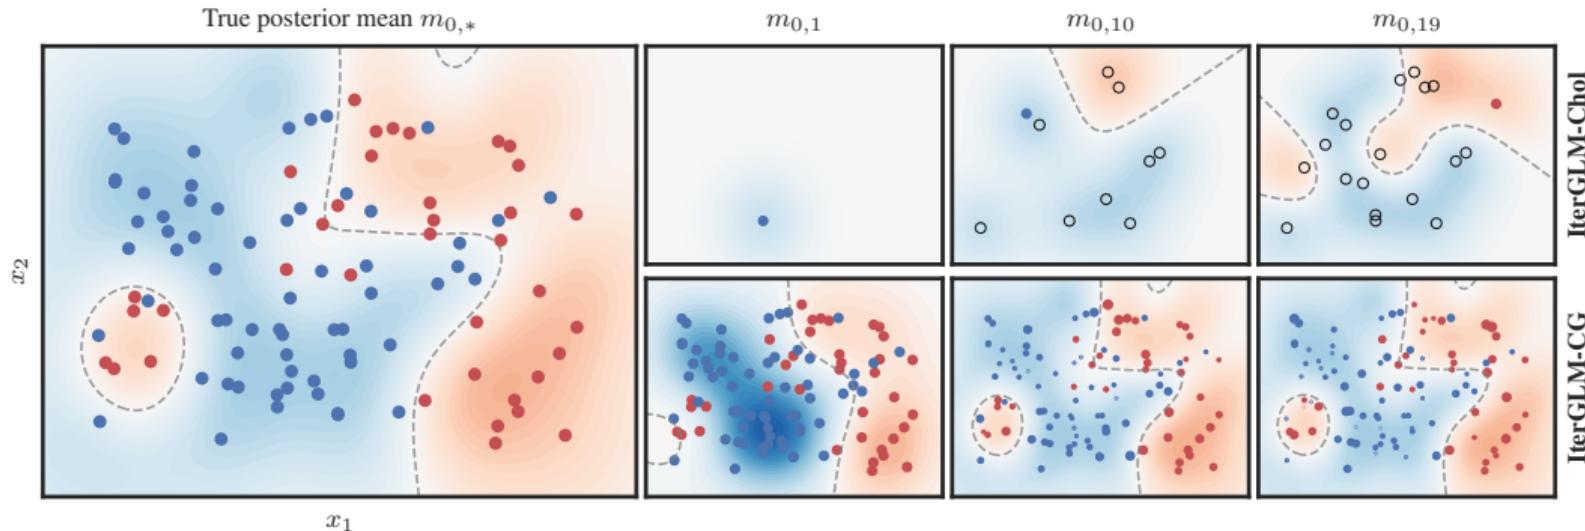




# Generalized Linear Models

Gaussian process classification with IterGLM using two different policies.

[Tat+23]





# Wrapping Up

## Summary

- ▶ Large-scale linear systems are ubiquitous in scientific computation.



# Wrapping Up

## Summary

- ▶ Large-scale linear systems are ubiquitous in scientific computation.
- ▶ Probabilistic linear solvers allow us to explicitly trade off **speed for precision**.



# Wrapping Up

## Summary

- ▶ Large-scale linear systems are ubiquitous in scientific computation.
- ▶ Probabilistic linear solvers allow us to explicitly trade off **speed for precision**.

## Application: Gaussian Processes

- ▶ Large-scale GP models are often as much about the approximation as they are about the data.



# Wrapping Up

## Summary

- ▶ Large-scale linear systems are ubiquitous in scientific computation.
- ▶ Probabilistic linear solvers allow us to explicitly trade off **speed for precision**.

## Application: Gaussian Processes

- ▶ Large-scale GP models are often as much about the approximation as they are about the data.



# Wrapping Up

## Summary

- ▶ Large-scale linear systems are ubiquitous in scientific computation.
- ▶ Probabilistic linear solvers allow us to explicitly trade off **speed** for **precision**.

## Application: Gaussian Processes

- ▶ Large-scale GP models are often as much about the approximation as they are about the data.
- ▶ We can exactly quantify the error from **finite data** and from the **approximation** via a combined uncertainty estimate.  $\Rightarrow$  IterGP



# Wrapping Up

## Summary

- ▶ Large-scale linear systems are ubiquitous in scientific computation.
- ▶ Probabilistic linear solvers allow us to explicitly trade off **speed for precision**.

## Application: Gaussian Processes

- ▶ Large-scale GP models are often as much about the approximation as they are about the data.
- ▶ We can exactly quantify the error from **finite data** and from the **approximation** via a combined uncertainty estimate.  $\Rightarrow$  IterGP
- ▶ **Explicit trade-off** between computation and uncertainty via probabilistic linear solver.



# Wrapping Up

## Summary

- ▶ Large-scale linear systems are ubiquitous in scientific computation.
- ▶ Probabilistic linear solvers allow us to explicitly trade off **speed for precision**.

## Application: Gaussian Processes

- ▶ Large-scale GP models are often as much about the approximation as they are about the data.
- ▶ We can exactly quantify the error from **finite data** and from the **approximation** via a combined uncertainty estimate.  $\Rightarrow$  IterGP
- ▶ **Explicit trade-off** between computation and uncertainty via probabilistic linear solver.

## Open Research Questions / Future Directions

- ▶ Calibration.
- ▶ Policy design for downstream tasks (Active learning, Bayesian optimization, ...).
- ▶ ...?



# Table of Contents

## 1 Introduction

## 2 Probabilistic Linear Solvers

2.1 Derivation

2.2 Policy Choice

2.3 Prior Choice

2.4 Algorithm

## 3 Application: Large-Scale Gaussian Processes

3.1 Gaussian Process Inference at Scale

3.2 Quantifying Approximation Error

3.3 Algorithm: IterGP

3.4 Theoretical Analysis

3.5 Policy Choice Illustrated

3.6 Experiments

## 4 Summary and Extensions

## 5 Additional Material

5.1 Calibration

5.2 An approximation method or a better model?

# Additional Material



# Calibration of BayesCG

Why is uncertainty quantification sometimes conservative for probabilistic linear solvers?

◀ ToC

**Observation:** Uncertainty quantification of probabilistic linear solvers can be conservative!

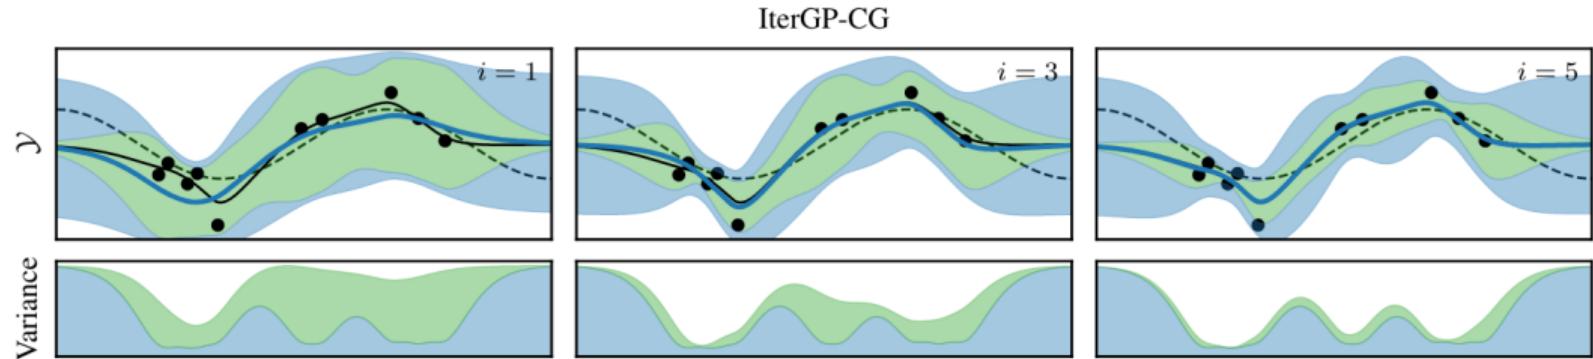


Figure: IterGP using a (conjugate) gradient policy.

# Calibration of BayesCG

Why is uncertainty quantification sometimes conservative for probabilistic linear solvers?

◀ ToC

**Observation:** Uncertainty quantification of probabilistic linear solvers can be conservative!

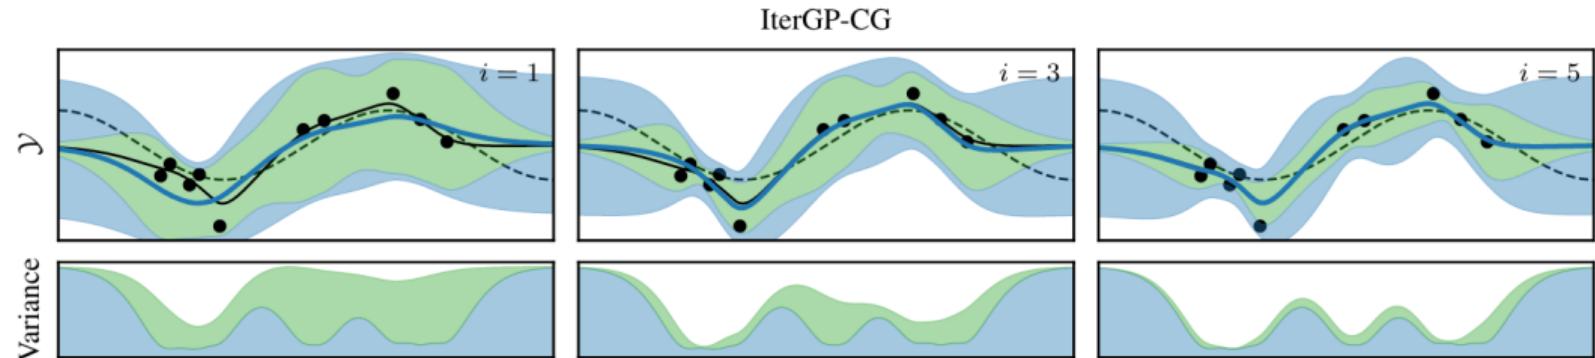


Figure: IterGP using a (conjugate) gradient policy.

**Why is that?** We conditioned on  $\alpha_i = \mathbf{s}_i^\top \mathbf{r}_{i-1} = \mathbf{s}_i^\top \mathbf{A}(\mathbf{x}_* - \mathbf{x}_{i-1})$ .

**But:** We've “cheated” for a gradient policy, since  $\mathbf{s}_i = \mathbf{b} - \mathbf{A}\mathbf{x}_{i-1} = \mathbf{A}(\mathbf{x}_* - \mathbf{x}_{i-1}) = \mathbf{s}_i(\mathbf{x}_*)$ .



# Working with Infinite Data

For IterGP it does not matter how large the dataset is, or whether we have it stored on our machine.

[Wen+22]

◀ ToC

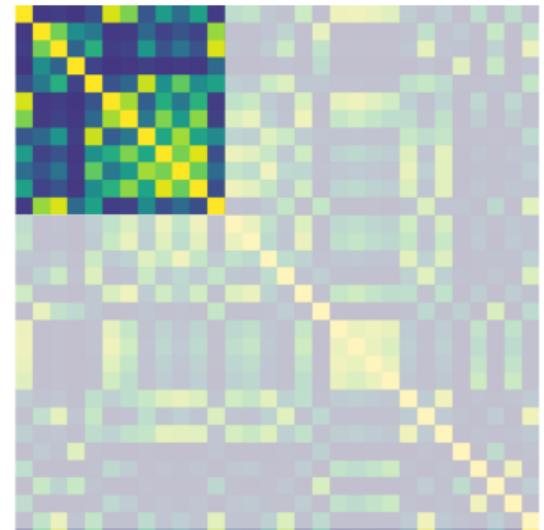
## Theorem (Online GP Approximation with IterGP)

Let  $n, n' \in \mathbb{N}$  and consider training data sets  $\mathbf{X} \in \mathbb{R}^{n \times d}, \mathbf{y} \in \mathbb{R}^n$  and  $\mathbf{X}' \in \mathbb{R}^{n' \times d}, \mathbf{y}' \in \mathbb{R}^{n'}$ . Consider two sequences of actions  $(\mathbf{s}_i)_{i=1}^n \in \mathbb{R}^n$  and  $(\tilde{\mathbf{s}}_i)_{i=1}^{n+n'} \in \mathbb{R}^{n+n'}$  such that

$$\tilde{\mathbf{s}}_i = \begin{pmatrix} \mathbf{s}_i \\ \mathbf{0} \end{pmatrix} \quad (1)$$

Then the posterior returned by IterGP for the dataset  $(\mathbf{X}, \mathbf{y})$  using actions  $\mathbf{s}_i$  is identical to the posterior returned by IterGP for the extended dataset using actions  $\tilde{\mathbf{s}}_i$ :

$$ITERGP(\mu, k, \mathbf{X}, \mathbf{y}, (\mathbf{s}_i)_i) = ITERGP \left( \mu, k, \begin{pmatrix} \mathbf{X} \\ \mathbf{X}' \end{pmatrix}, \begin{pmatrix} \mathbf{y} \\ \mathbf{y}' \end{pmatrix}, (\tilde{\mathbf{s}}_i)_i \right).$$





# An Approximation Method or a Better Model?

An alternative view of IterGP as a better model for the way we do inference instead of an approximation.

◀ ToC

**Observation:** Only once we perform computation on data, does it enter our prediction.





# An Approximation Method or a Better Model?

An alternative view of IterGP as a better model for the way we do inference instead of an approximation.

◀ ToC

**Observation:** Only once we perform computation on data, does it enter our prediction.



The distinction between data and computation vanishes.



# An Approximation Method or a Better Model?

An alternative view of IterGP as a better model for the way we do inference instead of an approximation.

◀ ToC

**Observation:** Only once we perform computation on data, does it enter our prediction.



The distinction between data and computation vanishes.

What if we modeled this situation with a Gaussian process?

$$f \sim \mathcal{GP}(\mu, k)$$

$$\tilde{y} | f(X) \sim \mathcal{N}(S_i^T f(X), \sigma^2 S_i^T S_i)$$

$$f | X, \tilde{y} \sim \mathcal{GP}(\mu_i, k_i)$$



# An Approximation Method or a Better Model?

An alternative view of IterGP as a better model for the way we do inference instead of an approximation.

◀ ToC

**Observation:** Only once we perform computation on data, does it enter our prediction.



The distinction between data and computation vanishes.

What if we modeled this situation with a Gaussian process?

$$f \sim \mathcal{GP}(\mu, k)$$

$$\tilde{y} | f(X) \sim \mathcal{N}(S_i^T f(X), \sigma^2 S_i^T S_i)$$

$$f | X, \tilde{y} \sim \mathcal{GP}(\mu_i, k_i)$$

IterGP's combined posterior is equivalent to exact GP regression for linearly projected data.



# References I

- ▶ E. Daxberger, A. Kristiadi, A. Immer, R. Eschenhagen, M. Bauer, and P. Hennig. *Laplace Redux – Effortless Bayesian Deep Learning*. 2022. doi: [10.48550/arXiv.2106.14806](https://doi.org/10.48550/arXiv.2106.14806). URL: <http://arxiv.org/abs/2106.14806> (cit. on p. 7).
- ▶ P. Hennig, M. A. Osborne, and M. Girolami. “Probabilistic numerics and uncertainty in computations”. In: *Proceedings of the Royal Society of London A: Mathematical, Physical and Engineering Sciences* 471.2179 (2015) (cit. on pp. 9–12).
- ▶ J. Cockayne, C. J. Oates, T. J. Sullivan, and M. Girolami. “Bayesian Probabilistic Numerical Methods”. In: *SIAM Review* 61.4 (2019), pp. 756–789. doi: [10.1137/17M1139357](https://doi.org/10.1137/17M1139357) (cit. on pp. 9–12).
- ▶ P. Hennig, M. A. Osborne, and H. P. Kersting. *Probabilistic Numerics: Computation as Machine Learning*. CUP, 2022. ISBN: 978-1-316-68141-1. doi: [10.1017/9781316681411](https://doi.org/10.1017/9781316681411) (cit. on pp. 9–12).
- ▶ P. Hennig. “Probabilistic Interpretation of Linear Solvers”. In: *SIAM Journal on Optimization* 25.1 (2015), pp. 234–260 (cit. on pp. 9–21).



## References II

- ▶ J. Cockayne, C. J. Oates, I. C. Ipsen, and M. Girolami. "A Bayesian Conjugate Gradient Method (with Discussion)". In: *Bayesian Analysis* 14.3 (2019), pp. 937–1012. doi: [10.1214/19-BA1145](https://doi.org/10.1214/19-BA1145) (cit. on pp. 9–23, 28, 30, 31).
- ▶ J. Wenger and P. Hennig. "Probabilistic Linear Solvers for Machine Learning". In: *Advances in Neural Information Processing Systems (NeurIPS)*. 2020 (cit. on pp. 9–12).
- ▶ J. Wenger, G. Pleiss, M. Pförtner, P. Hennig, and J. P. Cunningham. "Posterior and Computational Uncertainty in Gaussian Processes". In: *Advances in Neural Information Processing Systems (NeurIPS)*. 2022 (cit. on pp. 28, 48–59, 66–68, 102).
- ▶ L. N. Trefethen and D. Bau. *Numerical Linear Algebra*. Society for Industrial and Applied Mathematics (SIAM), 1997 (cit. on p. 29).
- ▶ M. Kanagawa, P. Hennig, D. Sejdinovic, and B. K. Sriperumbudur. *Gaussian processes and kernel methods: A review on connections and equivalences*. 2018. arXiv: [1807.02582](https://arxiv.org/abs/1807.02582) (cit. on pp. 66–68).



## References III

- ▶ M. Titsias. "Variational learning of inducing variables in sparse Gaussian processes". In: *International Conference on Artificial Intelligence and Statistics (AISTATS)*. 2009 (cit. on pp. 81–84).
- ▶ J. Hensman, N. Fusi, and N. D. Lawrence. "Gaussian processes for big data". In: *Conference on Uncertainty in Artificial Intelligence (UAI)*. 2013 (cit. on pp. 81–84).
- ▶ M. Pförtner, I. Steinwart, P. Hennig, and J. Wenger. *Physics-Informed Gaussian Process Regression Generalizes Linear PDE Solvers*. 2023. DOI: [10.48550/arXiv.2212.12474](https://doi.org/10.48550/arXiv.2212.12474). URL: <http://arxiv.org/abs/2212.12474> (cit. on p. 89).
- ▶ L. Tatzel, J. Wenger, F. Schneider, and P. Hennig. *Accelerating Generalized Linear Models by Trading off Computation for Uncertainty*. 2023. DOI: [10.48550/arXiv.2310.20285](https://doi.org/10.48550/arXiv.2310.20285). URL: <http://arxiv.org/abs/2310.20285> (cit. on p. 90).